

Auto-scheduling Deep Neural Networks for Hardware

Xuan Yang, Mingyu Gao, Qiaoyi Liu, Ankita Nayak, Jing Pu, Jeff Setter, Steven Bell, Heonjae Ha, Priyanka Raina, Christos Kozyrakis, Mark Horowitz

Introduction

Motivation

- DNNs are widely used
- Various DNN accelerators were proposed
- Large parallelism and locality in DNNs

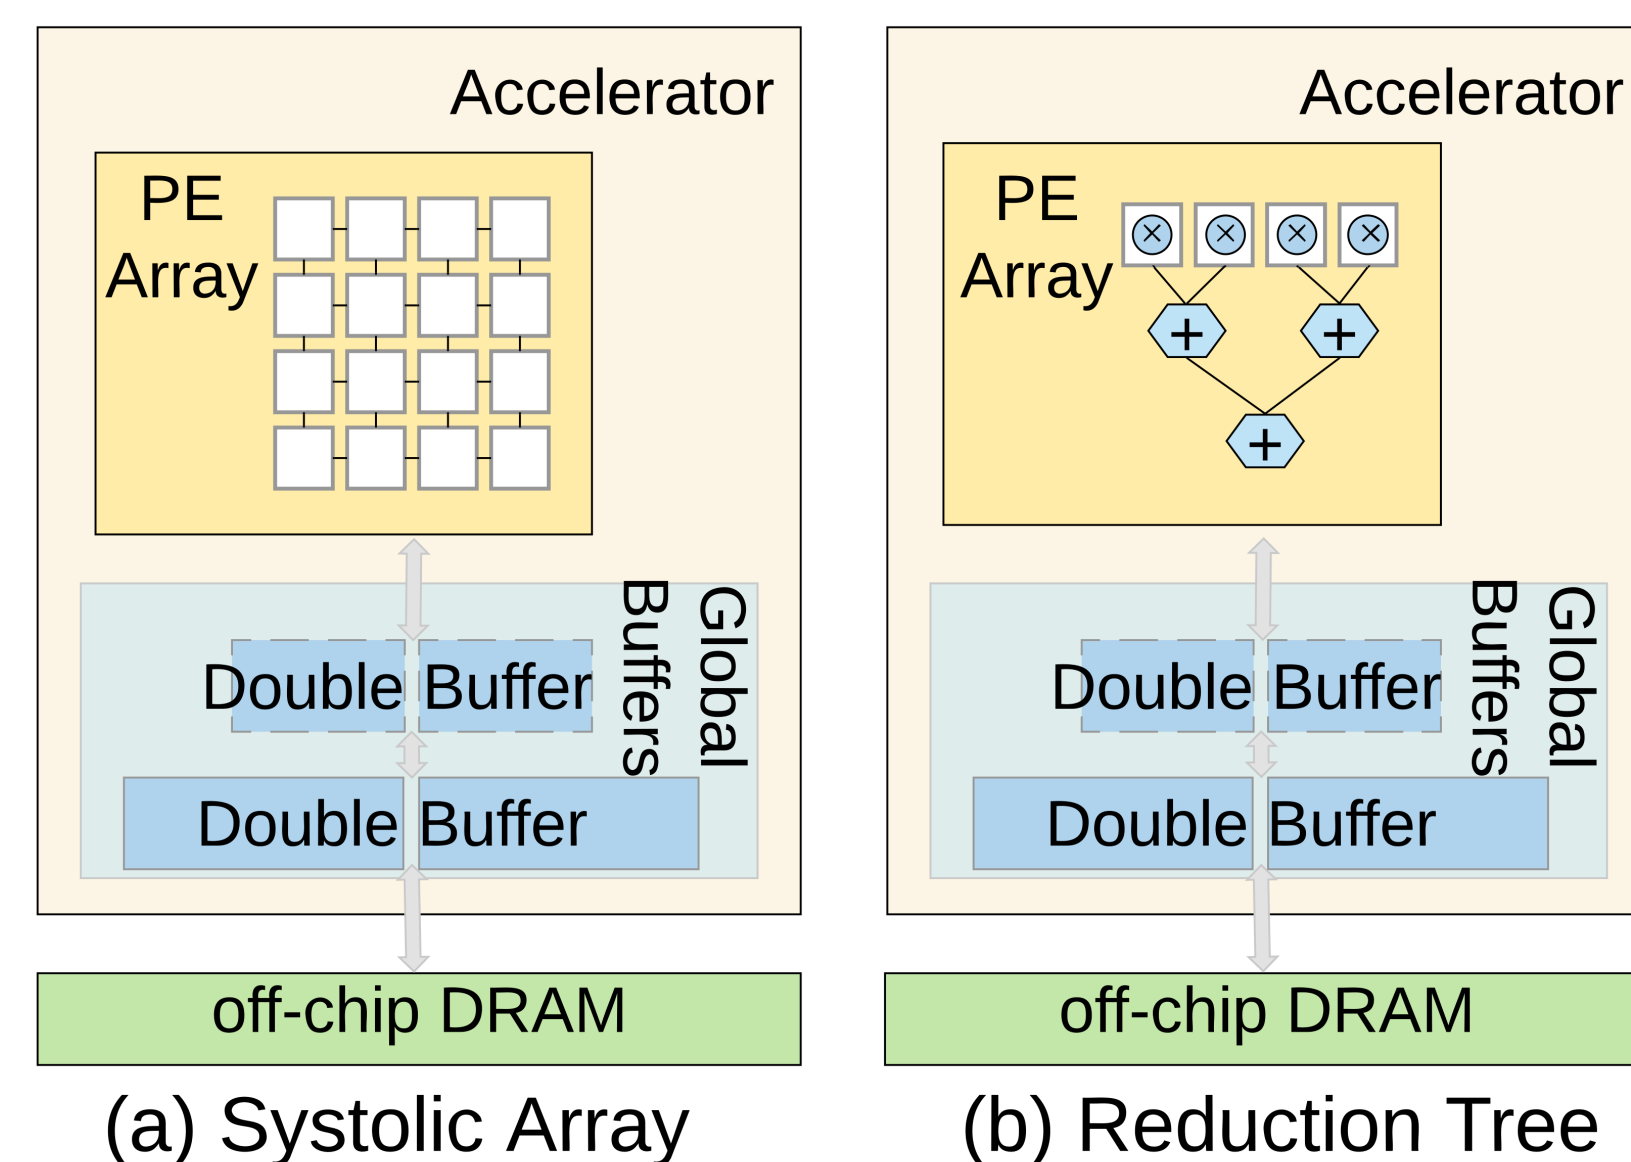
Architectural Optimization

- Balance different types of data reuse
- Loop blocking, parallelization


Language: Halide

- Halide splits *algorithm* from *schedule*
- Compact schedule to express hardware micro-architecture and mappings

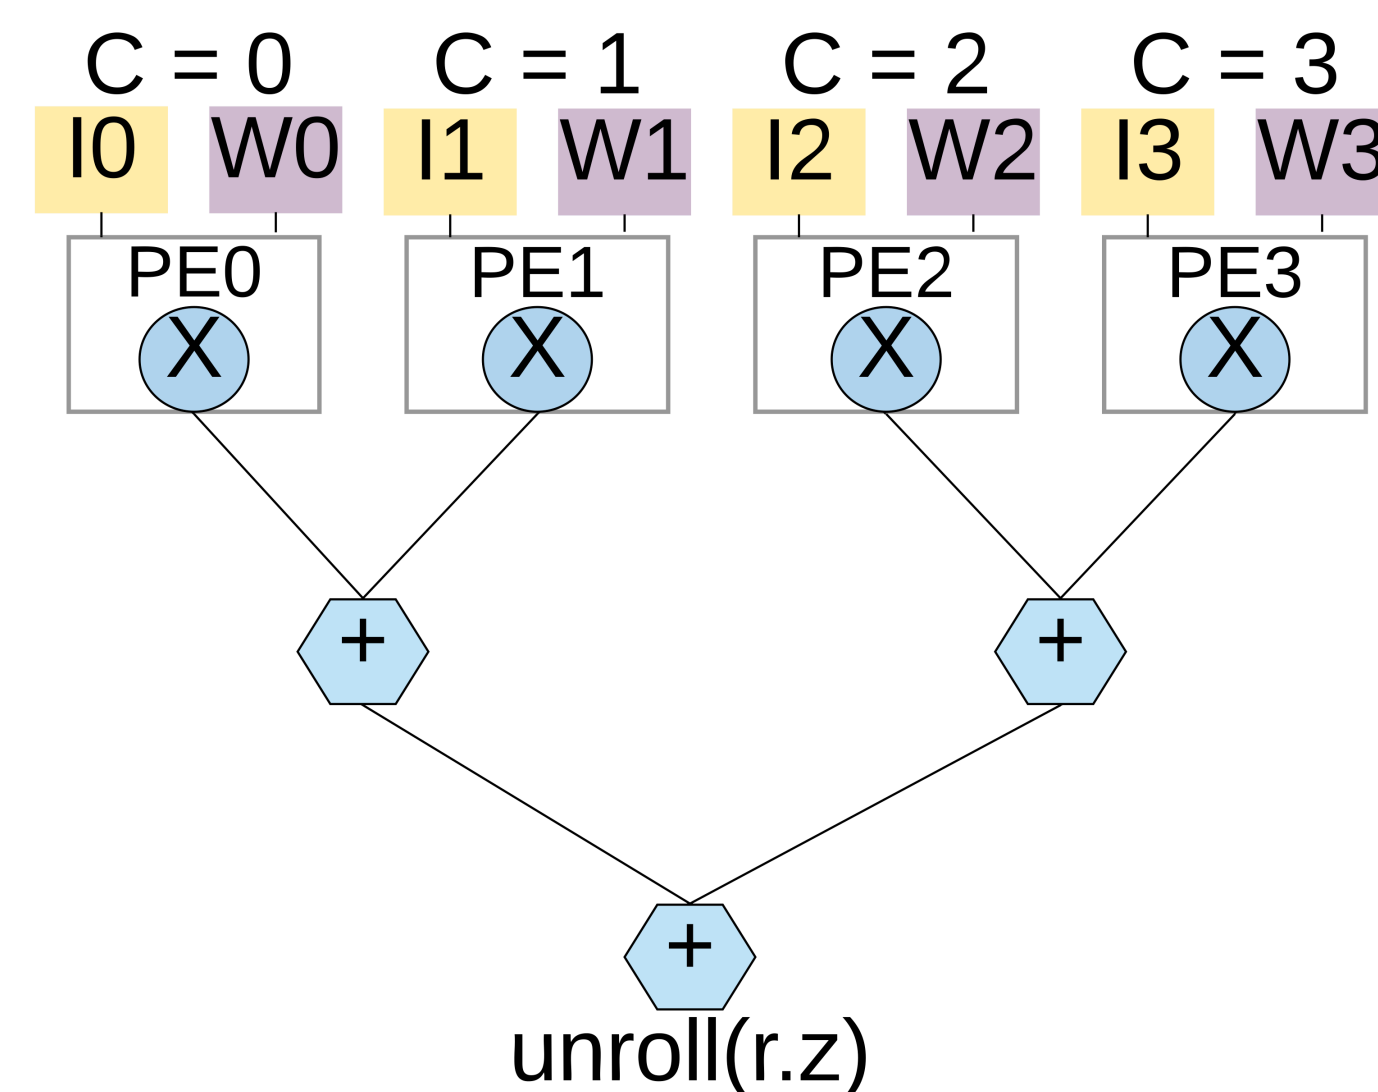
Architecture Template



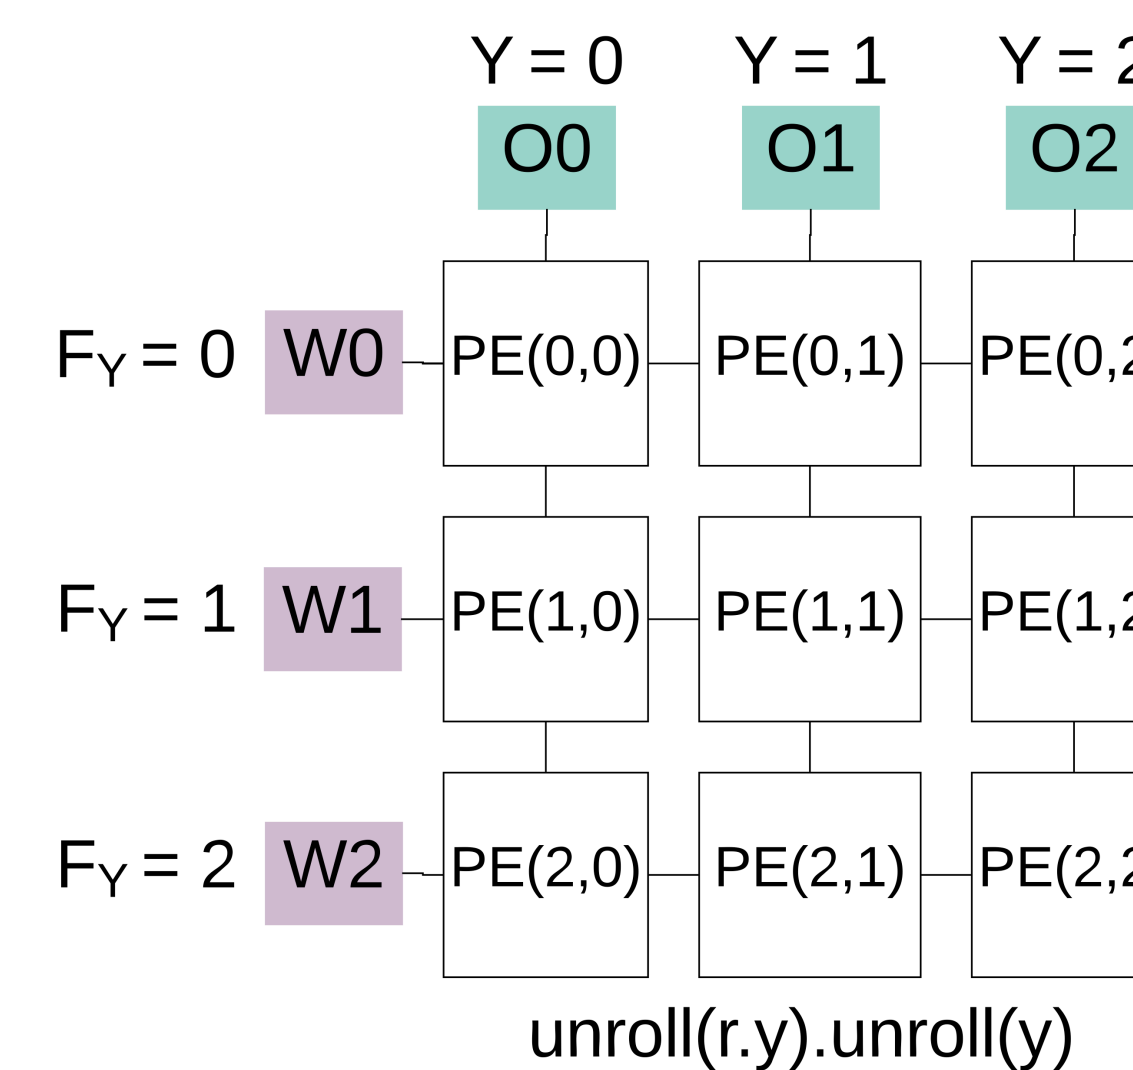
Architecture template for DNNs, the accelerator is composed of a PE array, a memory hierarchy with double buffers.



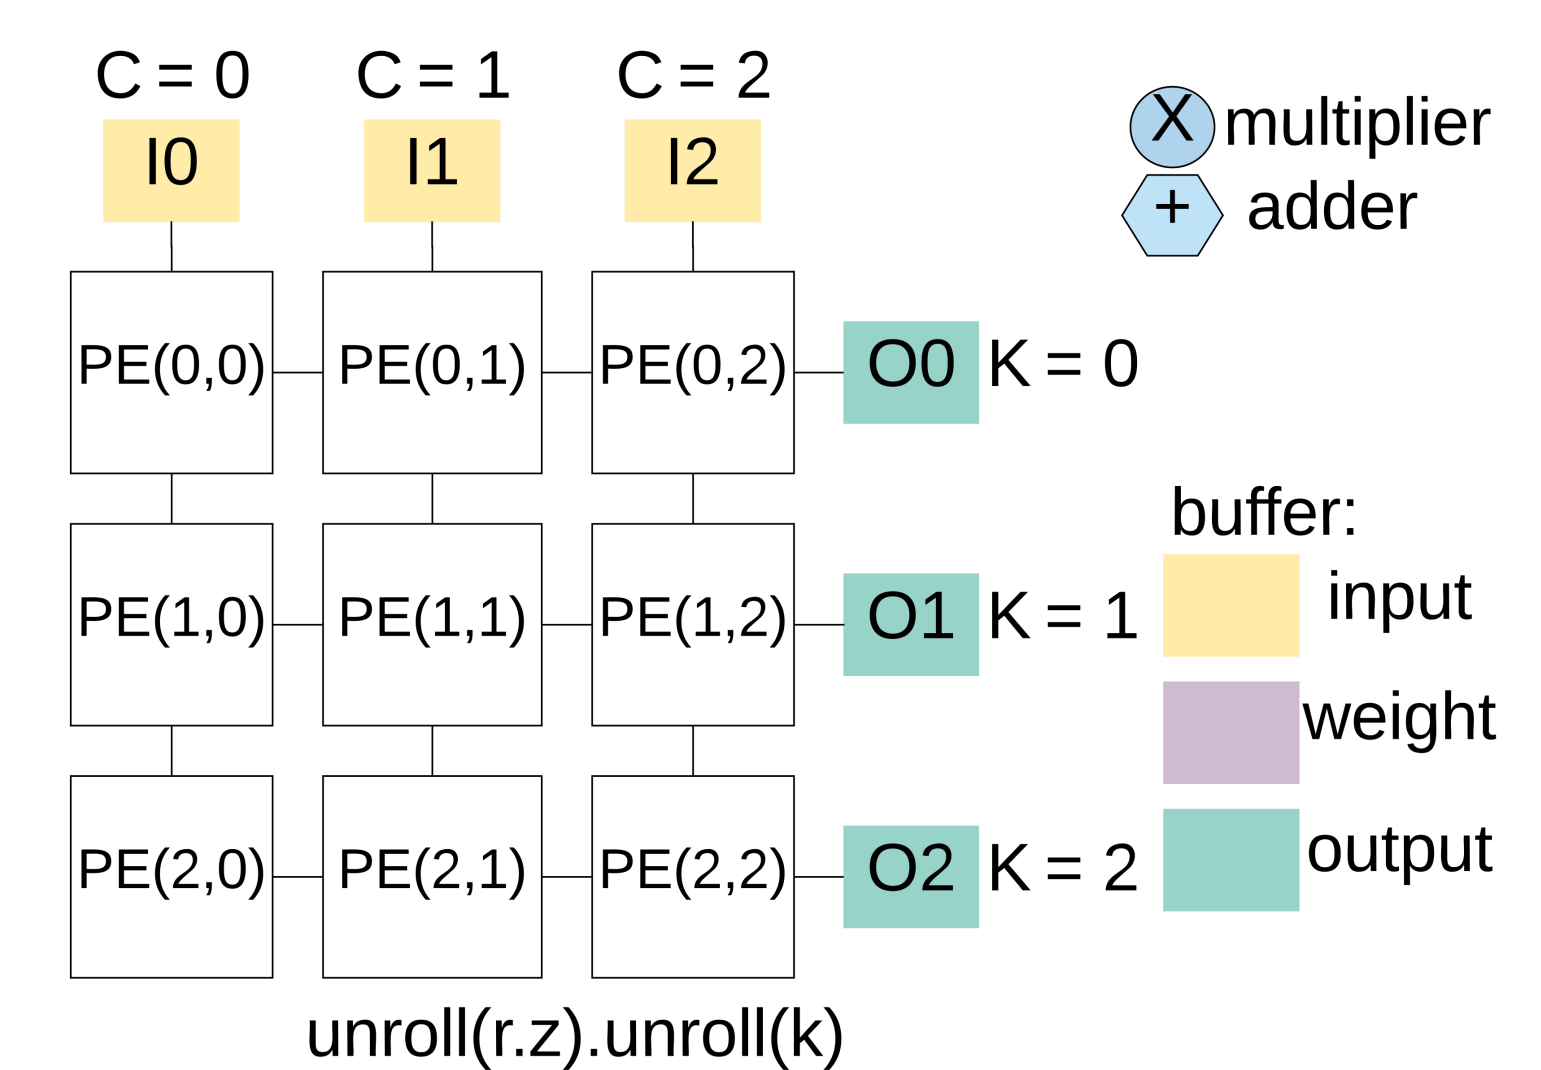
Micro-architectures Generation



Unrolling input channel dimension generates a reduction tree architecture.

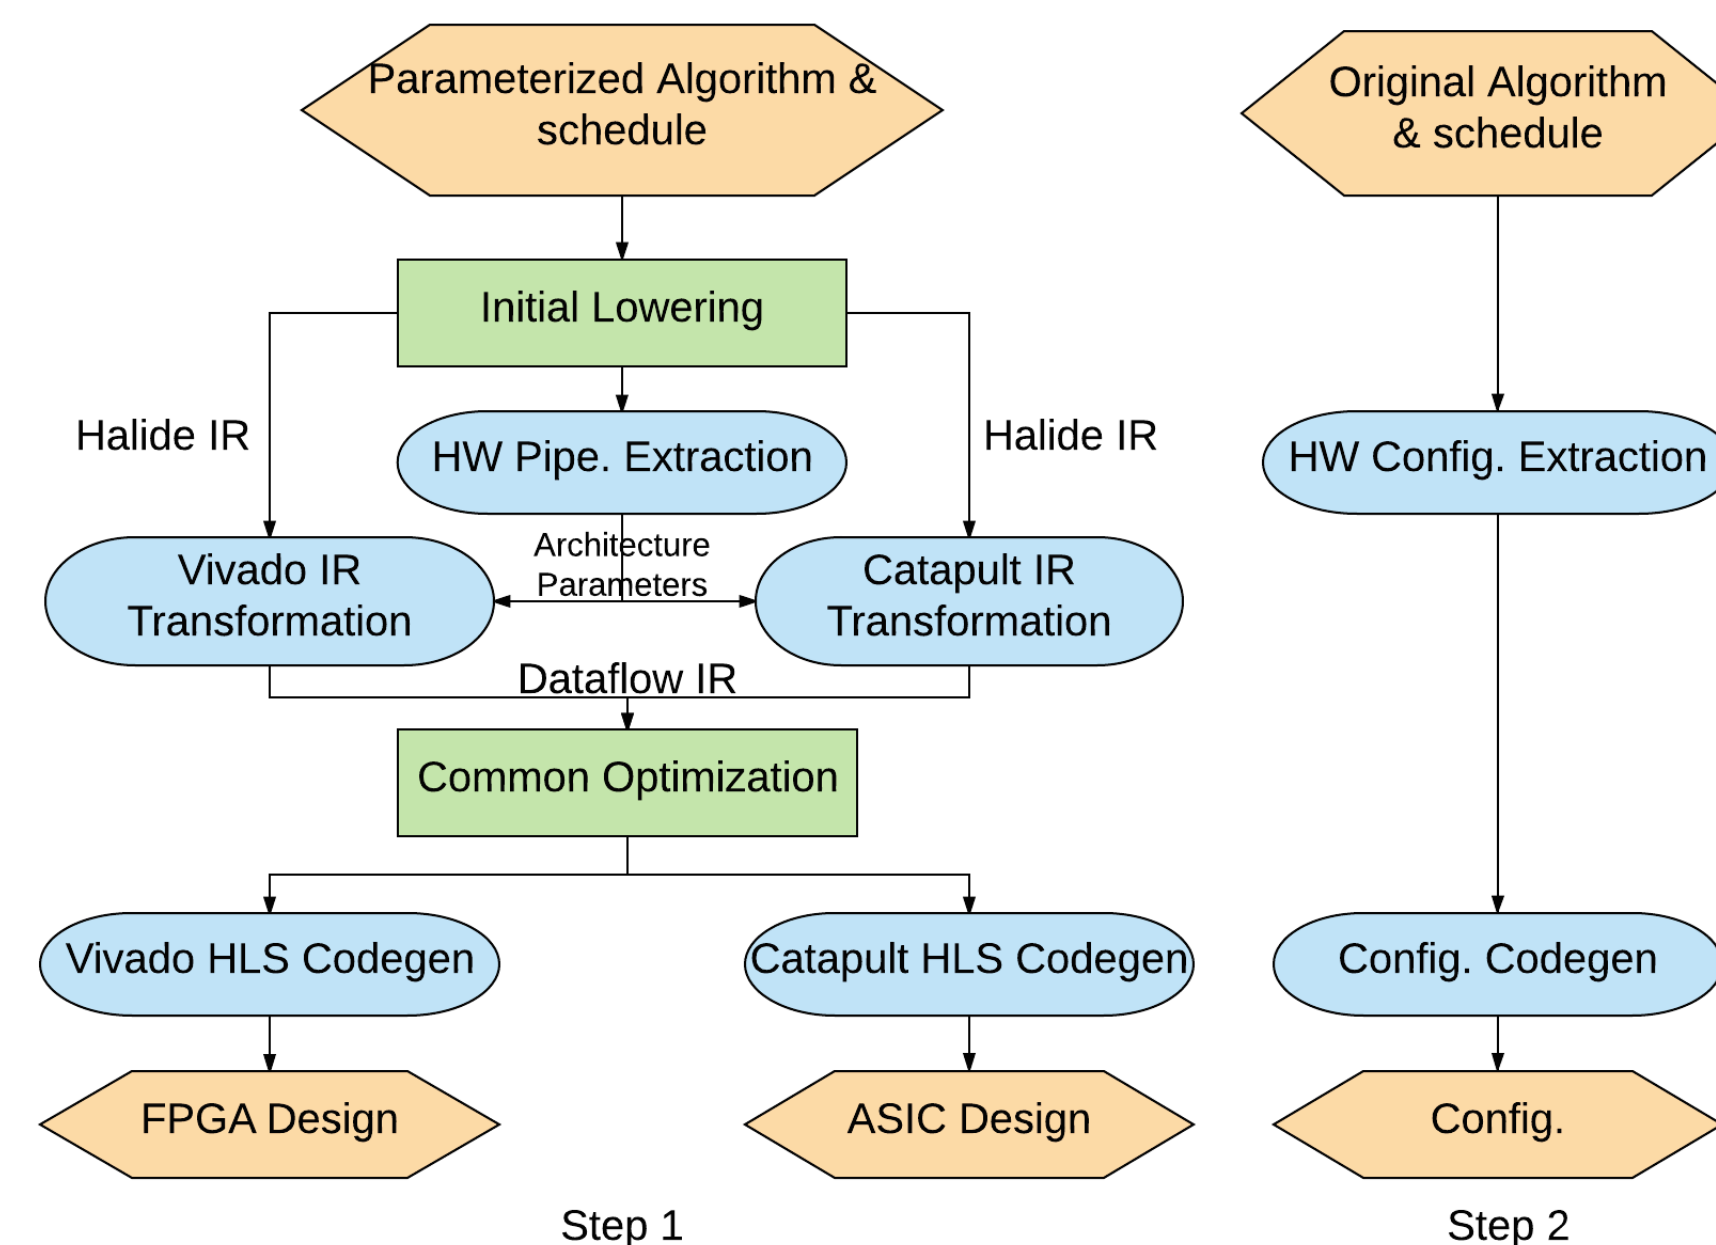


Unrolling image height and filter height dimensions, generates a systolic architecture, like Eyeriss.



Unrolling input channel and output channel dimensions, generates a systolic architecture, like TPU.


Virtualized Hardware Generation



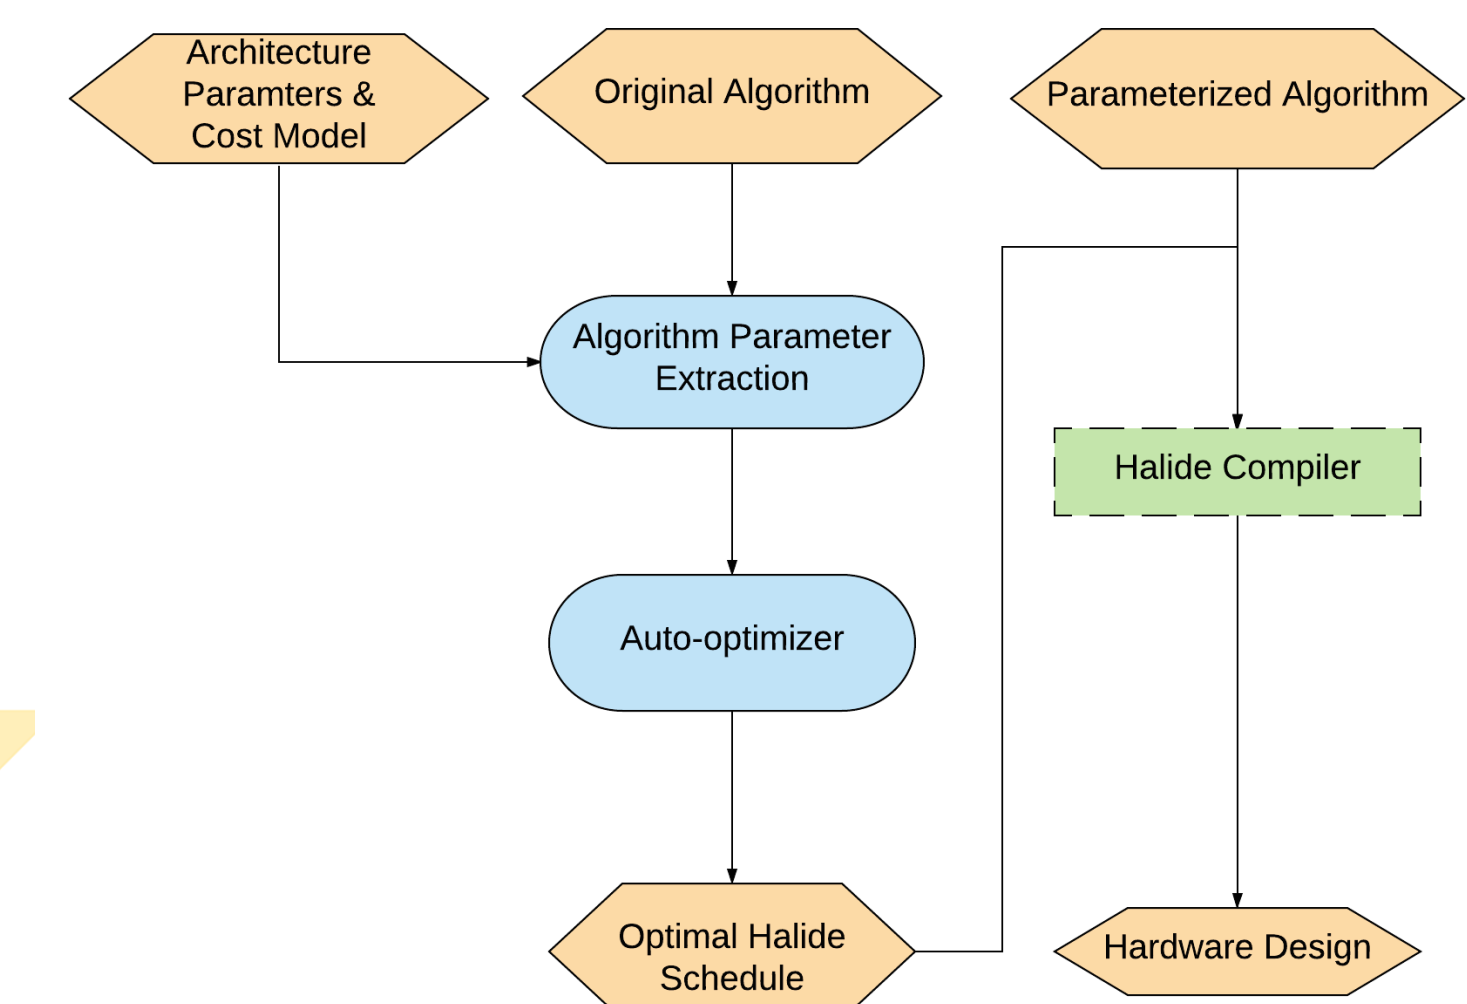
Two-step compilation flow. The first step is to generate a configurable hardware from a parameterized algorithm, the second step is to generate the configuration.



Auto-scheduler



The design space, Each hardware optimization dimension (*Blocking*, *Dataflow*, *HW Resource Allocation*) can be mapped to Halide schedule primitives (tile/reorder, unroll, in).



The auto-scheduler flow takes the original algorithm, architecture parameters and cost model to find the optimal schedule.

ISTC Agile