

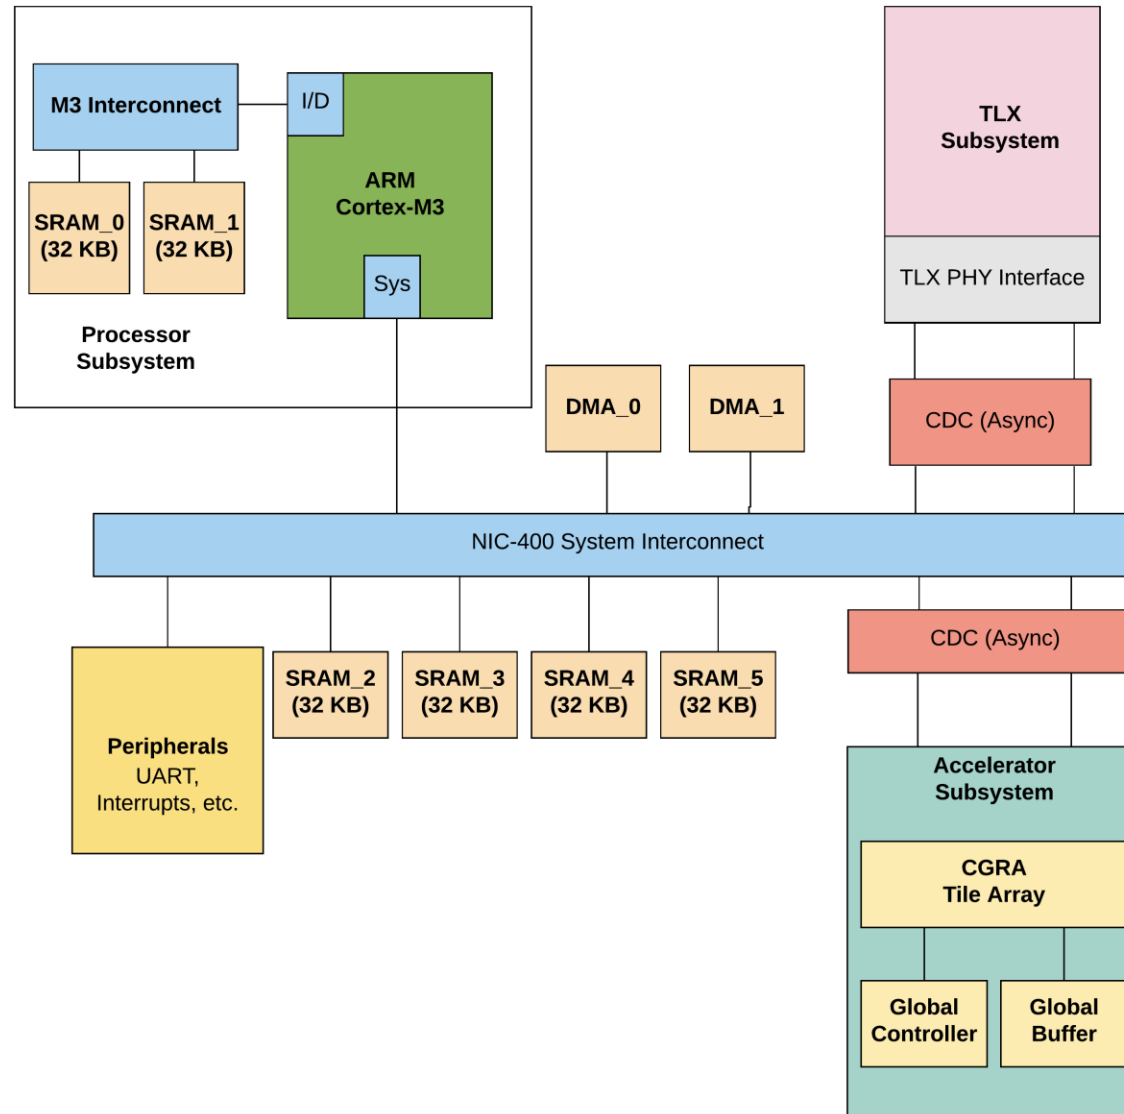
# CGRA Virtualization

Kalhan Koul

# Outline

- Amber SoC Review
  - Quick Overview of the SoC
  - Quick Overview of the CGRA Architecture
  - Application Walk Through
- Motivation for Virtualization
- Hardware/Software Enabling Improvements
- Temporal Virtualization
- Spatial Virtualization
- Ongoing Work

# Amber SoC Architecture



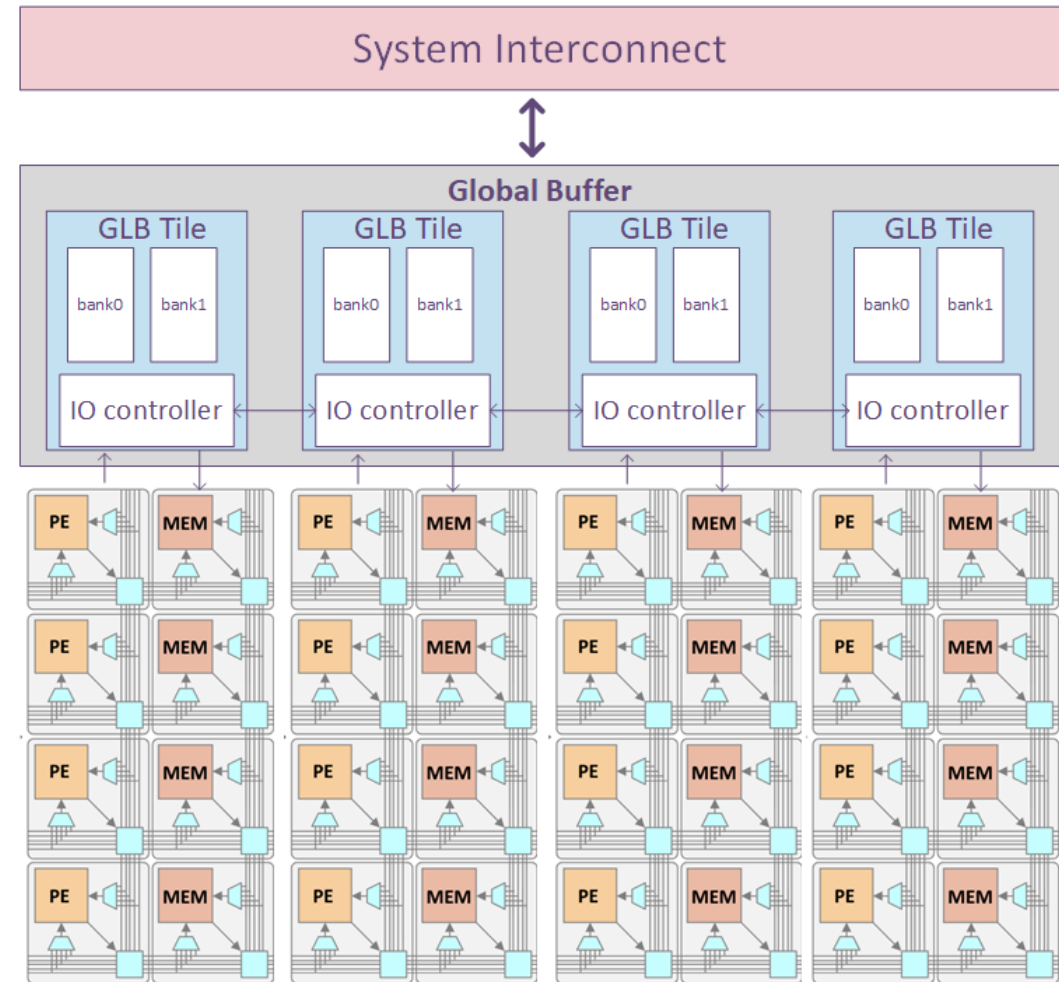
# CGRA Architecture

## Global Buffer

- Communicates with the SoC system over AXI4 and AXI4-Lite interface
- 1D array of GLB Tiles each with two banks
- I/O controller streams data into/out of CGRA
- Streams the bitstream into CGRA


## CGRA

- Island-style mesh network with PE tiles and MEM tiles
- Divided into groups of columns



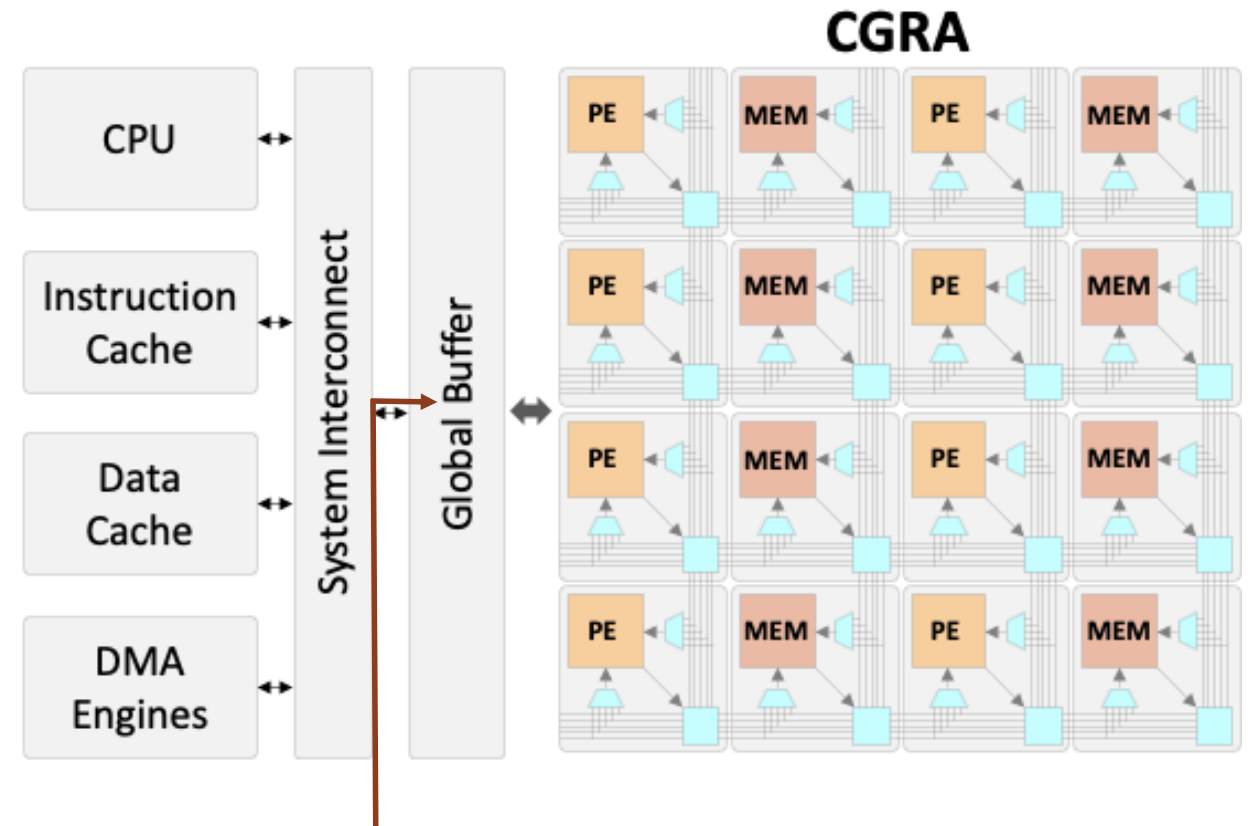
# Conducting an Application with the M3 Core

1. Initialize System
  - Enable UART, DMA engines, etc.
  - Select Clocks, Interrupts, etc.
2. Bring in Configuration and Input Data
  - Use DMA engines to move data (bitstream, images, etc.)
  - From DRAM/SRAMs to Global Buffer
3. Configure Global Buffer and Run Kernel
  - Use Parallel Configuration to send the bitstream to the CGRA
  - Stream in input data and stream out output data
4. Repeat Step 2 and 3 per Application
5. Finish Application
  - DMA Output Data to DRAM
  - Send done information to Host Processor



# Conducting an Application with the M3 Core


1. Initialize System
  - Enable UART, DMA engines, etc.
  - Select Clocks, Interrupts, etc.
2. Bring in Configuration and Input Data
  - Use DMA engines to move data (bitstream, images, etc.)
  - From DRAM/SRAMs to Global Buffer
3. Configure Global Buffer and Run Kernel
  - Use Parallel Configuration to send the bitstream to the CGRA
  - Stream in input data and stream out output data
4. Repeat Step 2 and 3 per Application
5. Finish Application
  - DMA Output Data to DRAM
  - Send done information to Host Processor



Off Chip: Bitstream, Input Data


# Conducting an Application with the M3 Core

1. Initialize System
  - Enable UART, DMA engines, etc.
  - Select Clocks, Interrupts, etc.
2. Bring in Configuration and Input Data
  - Use DMA engines to move data (bitstream, images, etc.)
  - From DRAM/SRAMs to Global Buffer
3. **Configure Global Buffer and Run Kernel**
  - **Use Parallel Configuration to send the bitstream to the CGRA**
  - **Stream in input data and stream out output data**
4. Repeat Step 2 and 3 per Application
5. Finish Application
  - DMA Output Data to DRAM
  - Send done information to Host Processor



# Conducting an Application with the M3 Core

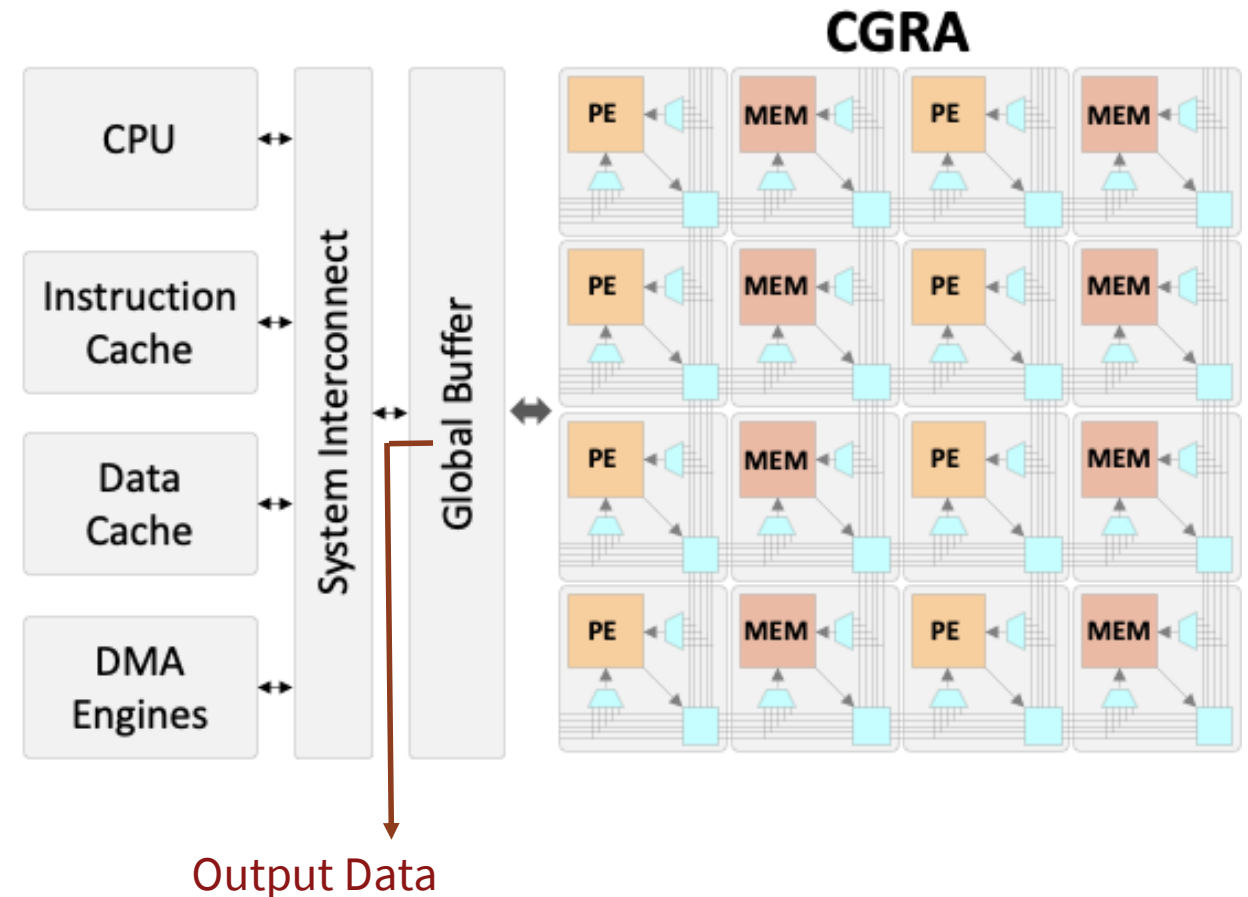
1. Initialize System
  - Enable UART, DMA engines, etc.
  - Select Clocks, Interrupts, etc.
2. Bring in Configuration and Input Data
  - Use DMA engines to move data (bitstream, images, etc.)
  - From DRAM/SRAMs to Global Buffer
3. Configure Global Buffer and Run Kernel
  - Use Parallel Configuration to send the bitstream to the CGRA
  - Stream in input data and stream out output data
4. **Repeat Step 2 and 3 per Application**
5. Finish Application
  - DMA Output Data to DRAM
  - Send done information to Host Processor






# Conducting an Application with the M3 Core

1. Initialize System
  - Enable UART, DMA engines, etc.
  - Select Clocks, Interrupts, etc.
2. Bring in Configuration and Input Data
  - Use DMA engines to move data (bitstream, images, etc.)
  - From DRAM/SRAMs to Global Buffer
3. Configure Global Buffer and Run Kernel
  - Use Parallel Configuration to send the bitstream to the CGRA
  - Stream in input data and stream out output data
4. Repeat Step 2 and 3 per Application
5. Finish Application
  - **DMA Output Data to DRAM**
  - **Send done information to Host Processor**

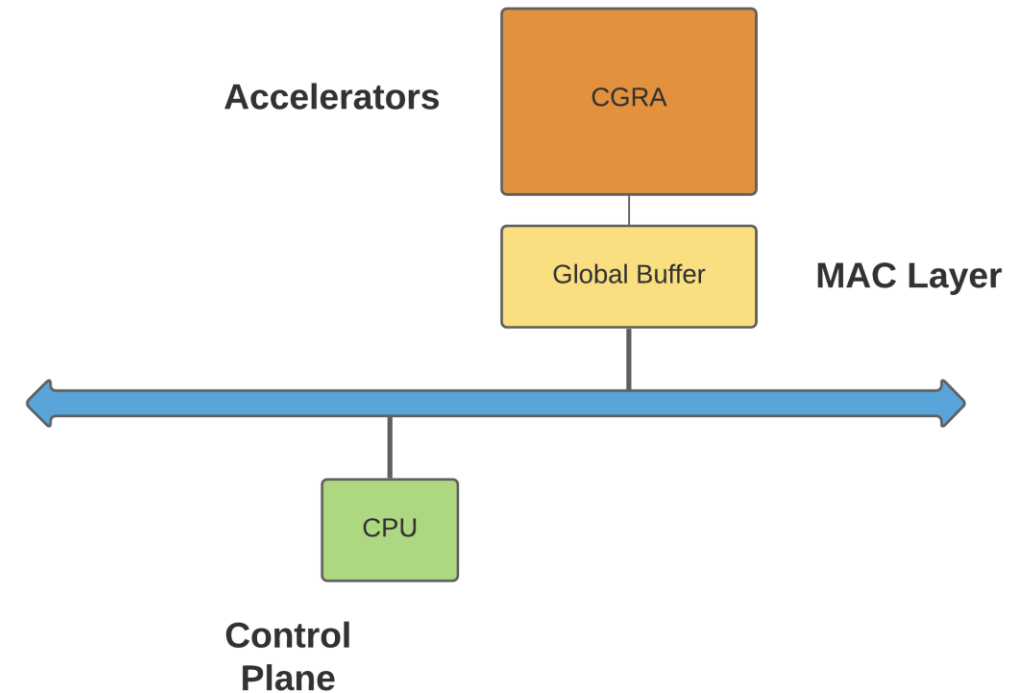


# Motivation for Virtualization

## Traditional View





## Our View




# Motivation for Virtualization

- 1. Cloud:** Two users access server for application acceleration
  - Most domain specific accelerators will schedule one after the other with high PE utilization
  - Large CGRA fabric gives us opportunity to do both in parallel, to achieve high PE utilization
- 2. Edge:** Two sensors need immediate processing
  - Most edge devices have several accelerators they need to manage and schedule
  - CGRA gives us opportunity to do both in parallel on the same accelerator, simplifying SoC architecture and scheduling management



# Hardware/Software Features Implemented

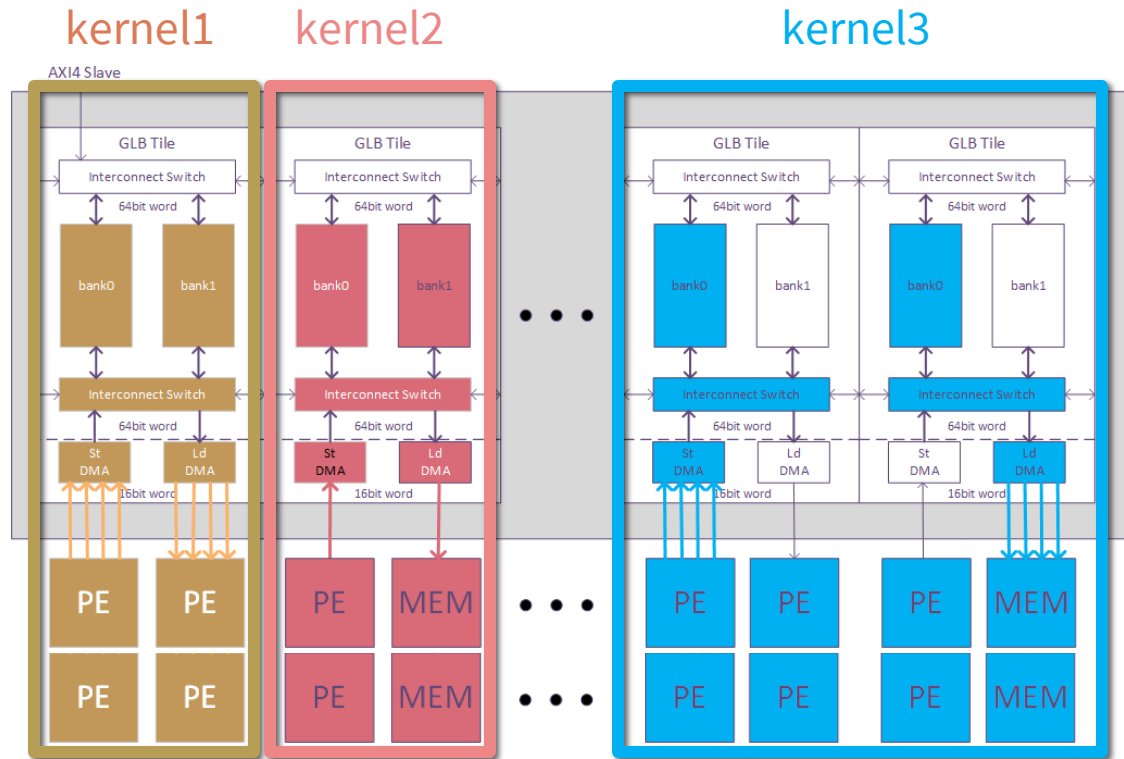
# Flexible Global Buffer Allocation



## For a Single Kernel

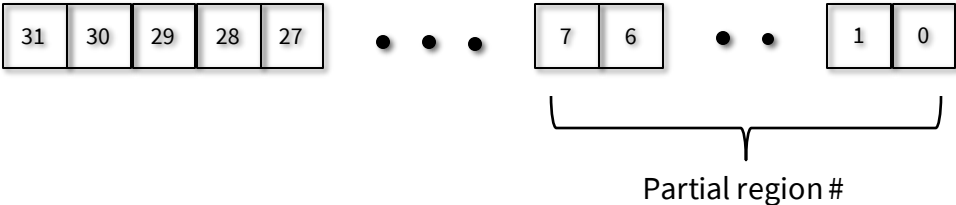
- Input buffer (4 banks)
- Output buffer (4 banks)
- Flexibility in mapping applications with different input/output data bandwidth

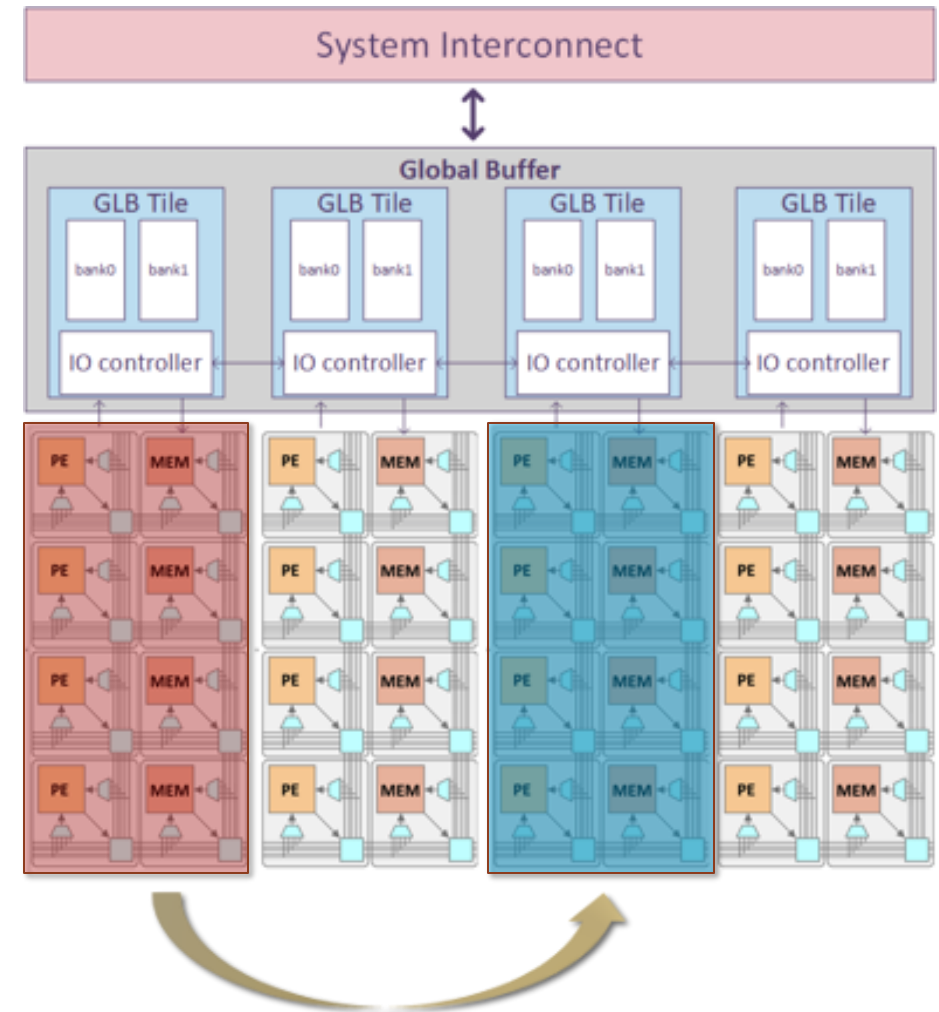
# Multiple Kernels in Parallel



- Up to 8 Kernels on 32x16 CGRA Fabric**
- Different kernels run in different regions of the CGRA
  - Independent control of each kernel (start, interrupt)

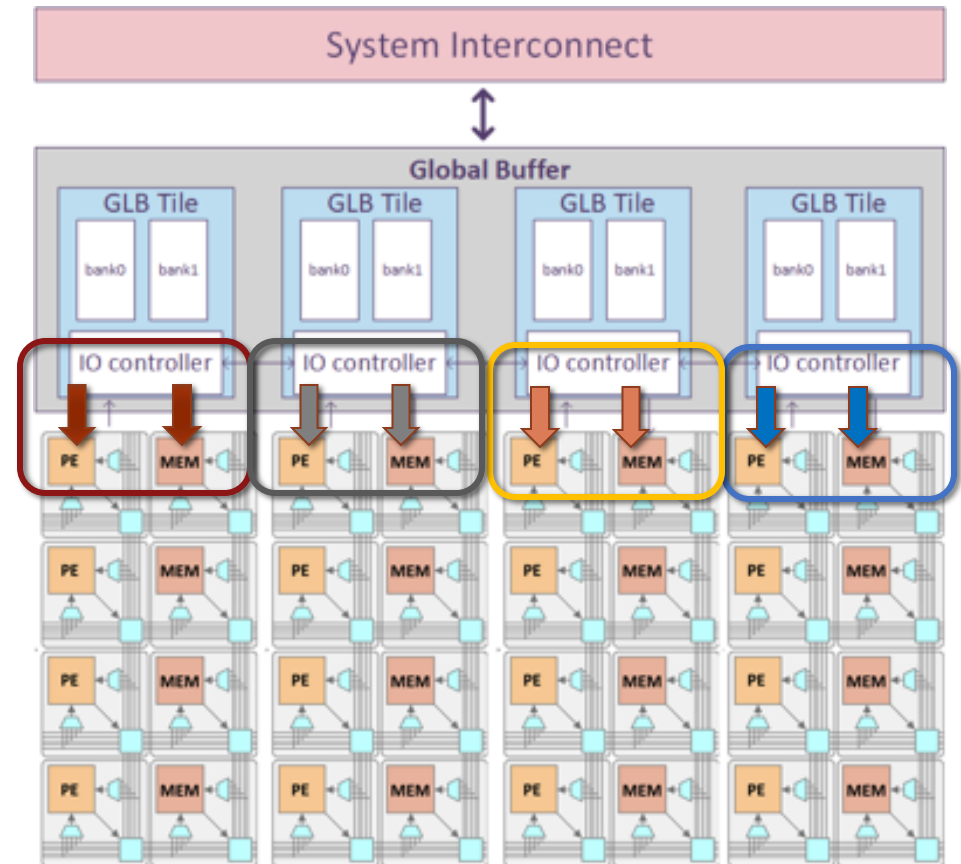
# Hardware Support for Reconfiguration

- #1 Bitstream is relocatable
  - › This allows flexible hardware allocation for each kernel without software overhead
- Bitstream address breakdown
  - › LSB 8bits indicate the partial region
  - › Rest of bitstream bits remain unchanged
- Bitstream address relocation
  - › Currently, this is done on the M3, which is slow
  - › Hardware to handle bitstream relocation will be ready soon








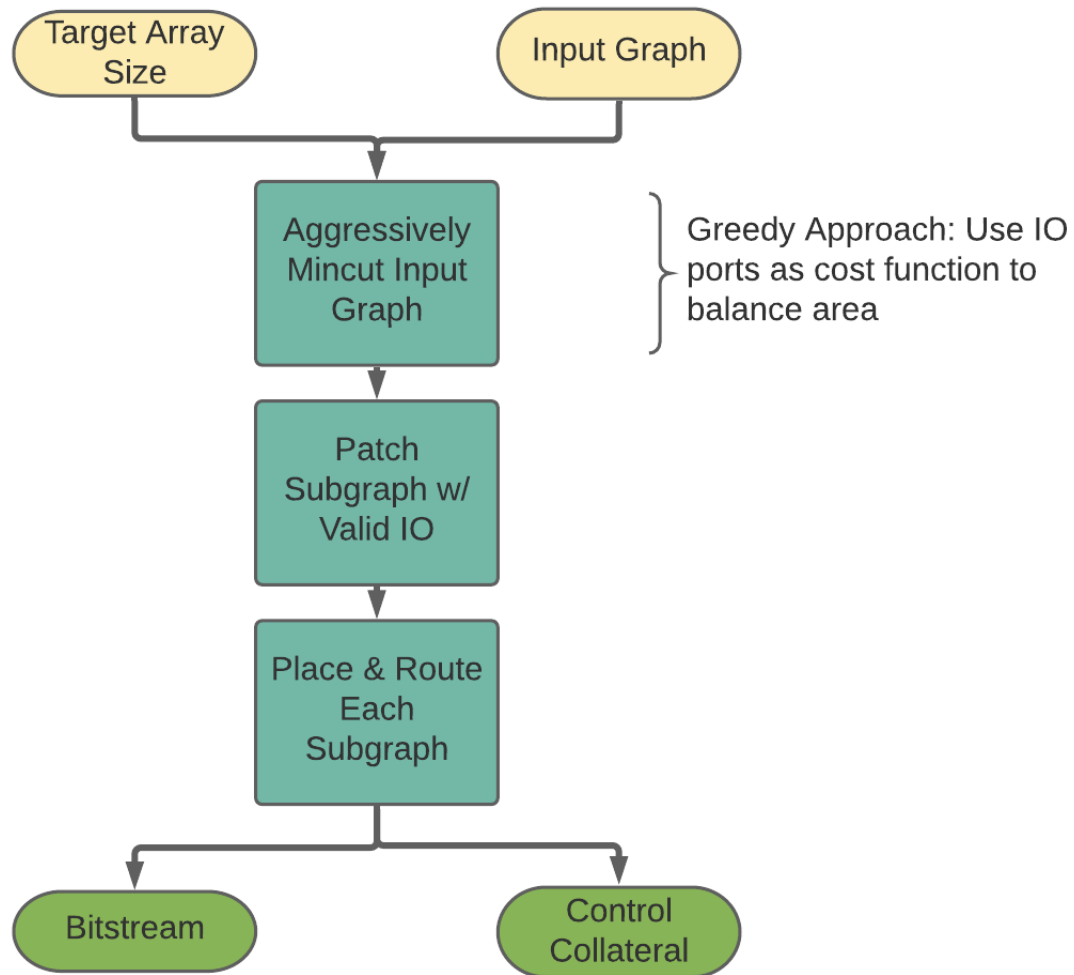
# Hardware Support for Dynamic Partial Reconfiguration

- #2 Dynamic Partial Reconfiguration (DPR) is fast, parallel, independent
  - › Bitstream is stored in global buffer tile
    - › Fast: Global buffer streams burst of bitstream
    - › Parallel: Up to 16 controllers can work in parallel
    - › Independent: Each controller is independent
- Configuration time (75,520 registers)
  - › JTAG: 108.7ms
  - › AXI4-Lite: 604us
  - › DPR: 4.72us





# Partitioning Algorithm

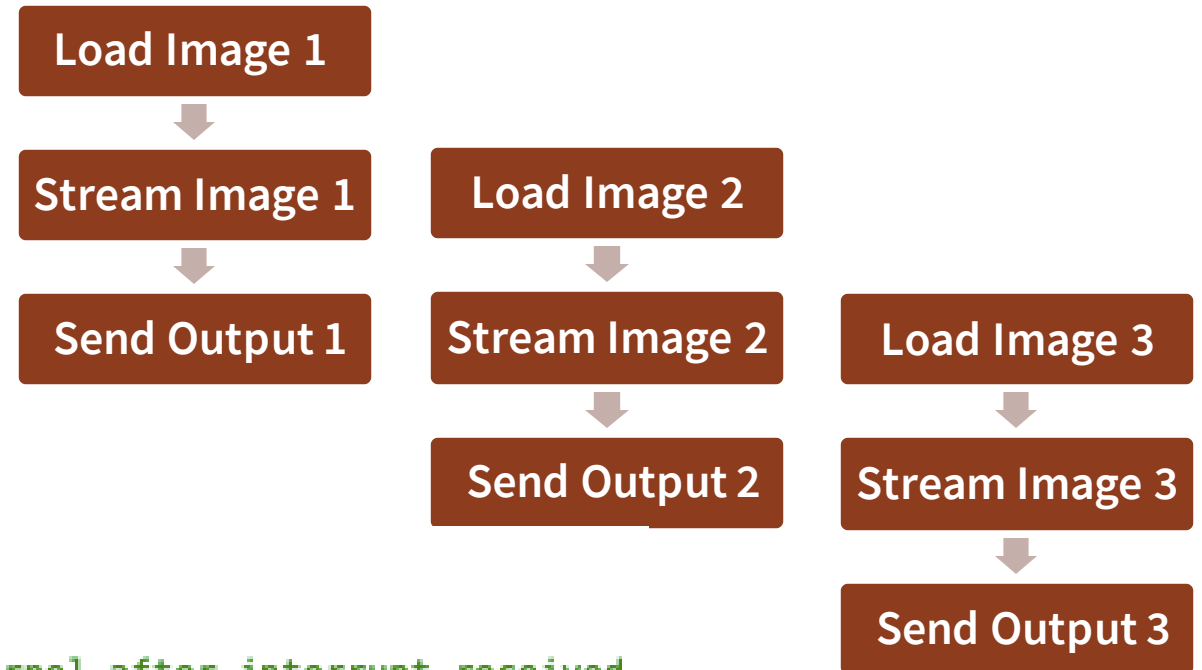


# Temporal and Spatial Virtualization

# Temporal Scheduling

- The M3 and Global Buffer pairing allow us to efficiently schedule applications
- For example: a batch of images with the same kernel
- Cycles to put an image in GB  $\approx$  cycles to stream an image to the fabric
- Can Send Output 1 and Load Image 3 at same time with 2 DMAs

```
// Pseudocode
// Time step 3
stream_image_to_fabric(2); // done with kernel after interrupt received
send_image_out_gb(1); // DMA1
load_image_to_gb(3); // DMA2
wait_for_interrupt();
```

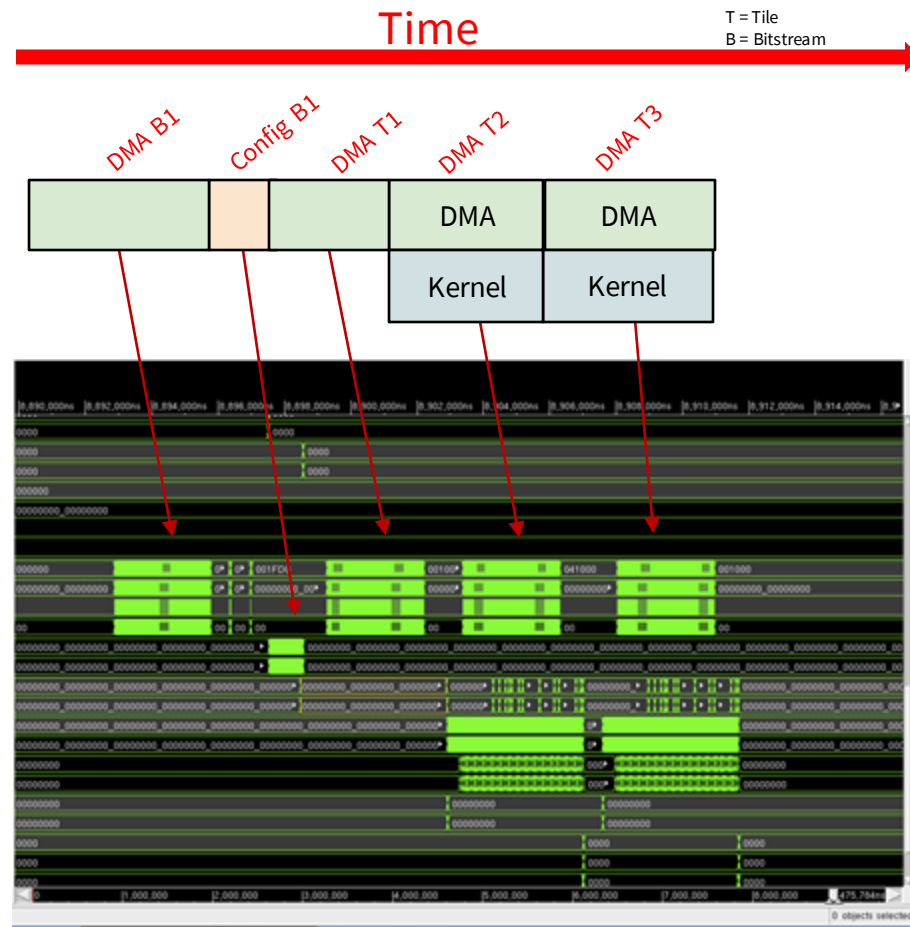


# Temporal Scheduling Results

## Scenario 1: Stream images into a single kernel (accelerator)

Result: **88% Utilization**

Metric	Description	Time
Programming Time	DMA bitstream and configure	6300 Cycles
Launch Time	Time between kernels	600 Cycles
Processing Time & Data Movement	Run kernel and DMA next tile	4096 Cycles

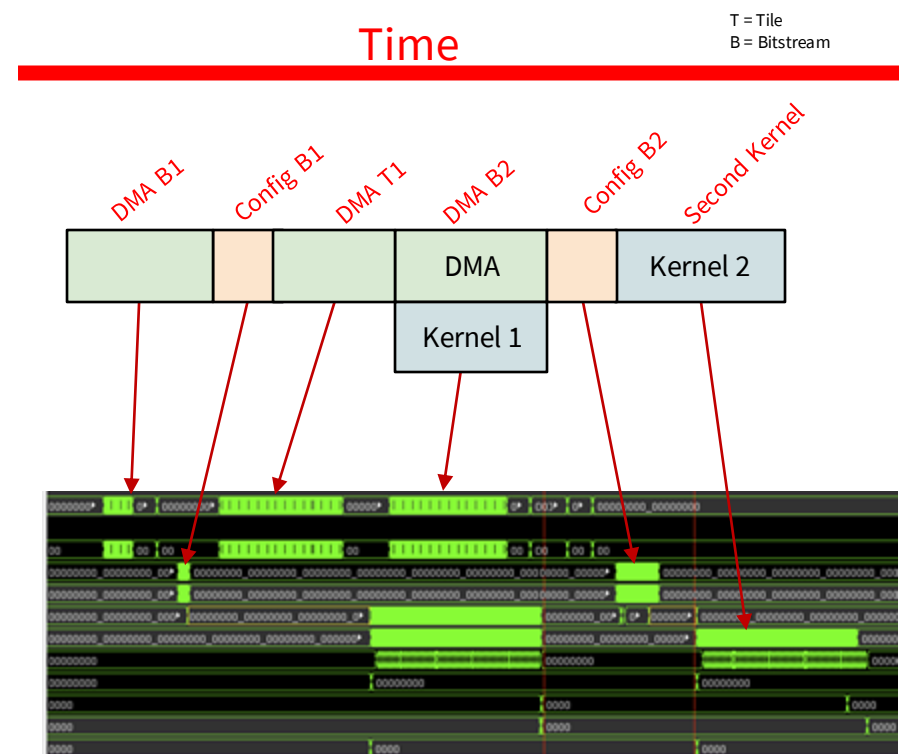


# Temporal Scheduling Results Cont.

## Scenario 2: Two different kernels back to back

Result: **70% Utilization**

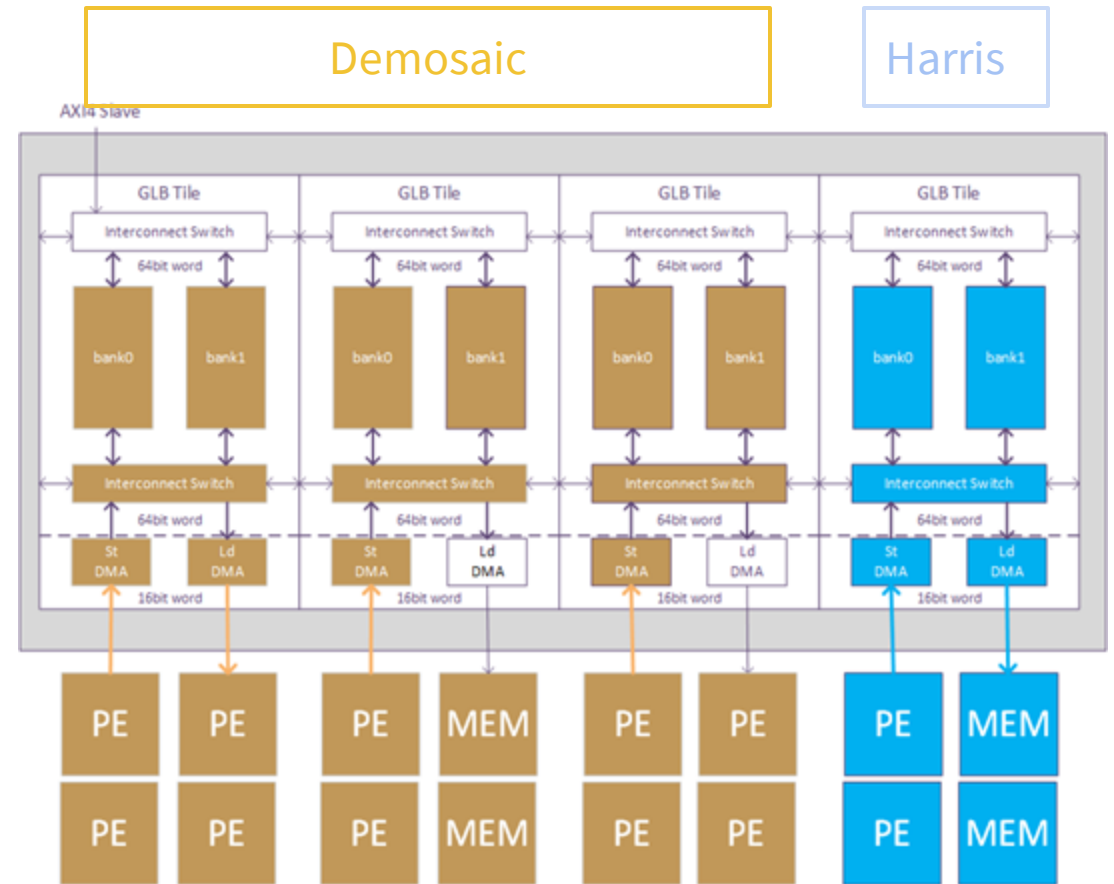
Metric	Description	Time
Programming Time	DMA bitstream and configure	6300 Cycles
Launch Time	Time between kernels	3500 Cycles
Processing Time & Data Movement	Run kernel and DMA next tile	4096 Cycles



# Spatial Scheduling


- We can place more than one kernel onto the fabric at once
- For example: Demosaic and Harris configured on the CGRA
- With 2 DMAs we can run both applications independently

```
// Pseudocode
// Time step 1
stream_config_to_fabric(1); // wfi
stream_config_to_fabric(2); // wfi
wait_for_interrupt(1 & 2);
```



# Spatial Scheduling Results

Individual App and Virtualized Power



# Characterization of Simultaneous Applications

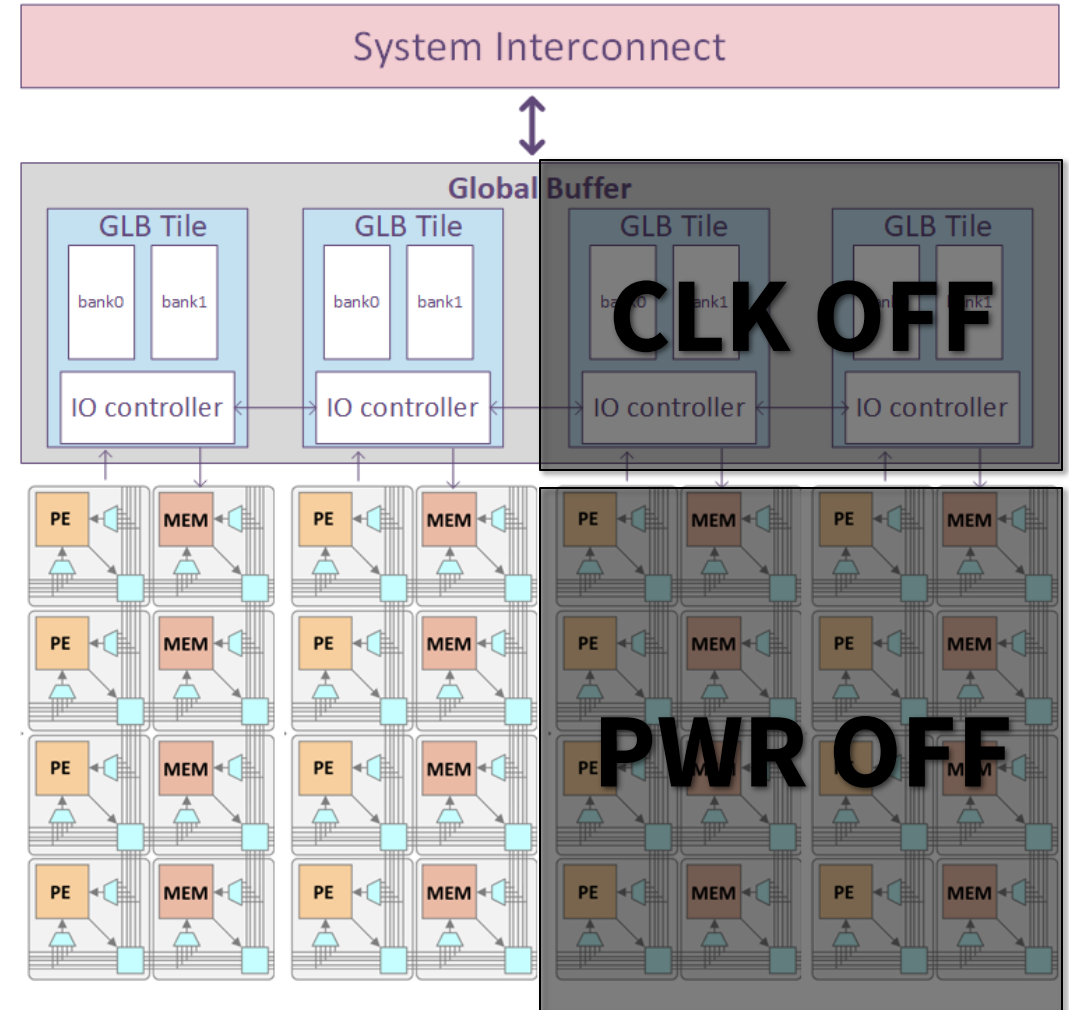
Metric	Demosaic only (unused tiles powered down)	Harris only (unused tiles powered down)	Harris and Demosaic simultaneously (unused tiles powered down)
Image size	480x640	480x640	480x640*2 Images
Tile size	64x64 (80 Tiles)	64x64 (80 Tiles)	64x64 (80 Tiles*2 Images)
Utilization (#PEs, #MEMs, #I/O, #Tracks)	217, 10, 6 16bit-track 14.4% 1bit-track 1.4%	83, 5, 4 16bit-track 11.5% 1bit-track 2.0%	300, 15, 10 16bit-track 13.7% 1bit-track 1.6%
Latency (startup latency, processing latency)	12000 Cycles, 6900 Cycles	6300 Cycle, 6900 Cycles	18300 Cycles, 6900 Cycles
Throughput @750MHz	2000 images/second	2000 images/second	4000 images/second
Energy	2.84pJ/op	7.19pJ/op	2.55pJ/op



# Ongoing Work

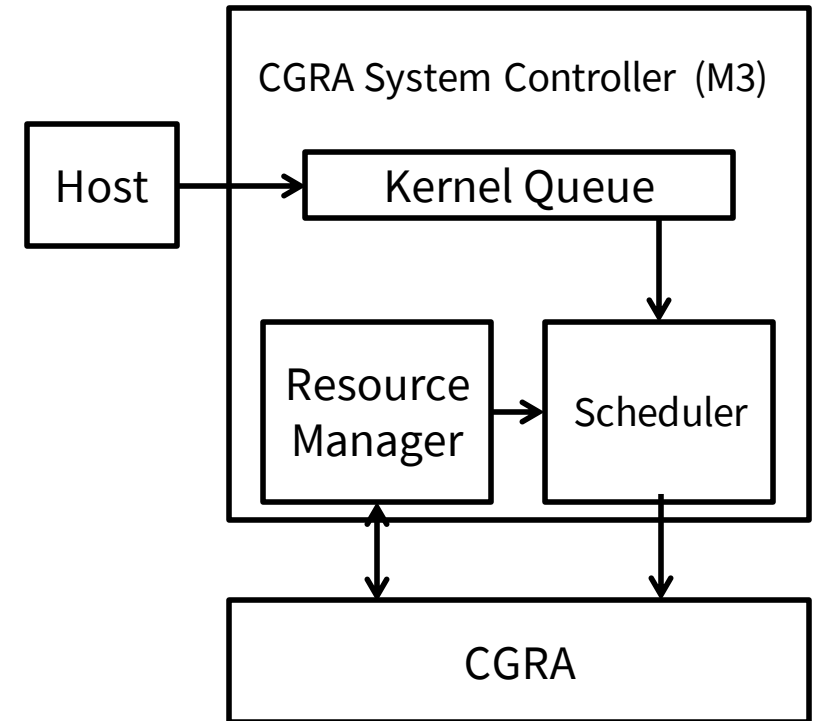
# Clock Gate Global Buffer Tiles

- Global buffer is a complex block with SRAMs, DMAs, and bitstream control logic
- Harris corner detection total power: 377mW
  - › Global buffer: 273mW
    - Used global buffer tiles: 41mW
    - Unused global buffer tiles: **232mW**
  - › Currently, Only SRAM macros are gated, but the idle global buffer tiles still consume a lot of power
- Future clock gating
  - › Power is mostly consumed by switching, so clock gating will save lots of power
  - › Configuration register controls clock-gating for each global buffer tile



# Dynamic scheduling

- CGRA System run-time controller on M3
  - Interface between host and CGRA
  - Controls kernels
- CGRA System controller components
  - Queue
    - Outstanding kernels are waiting here
  - Resource manager
    - Checks available hardware resources
  - Scheduler
    - Schedules kernels to hardware
    - Simple first-in-first-out scheduling
    - Complicated out-of-order scheduling



# Autoscheduling

- Leverage high degree of determinism on CGRA – utilize per operation and per memory access costs to create a configurable cost function based on performance, energy, and area
- Using an outer loop over different schedules and a cost function, we can determine the optimal schedule
- Can divide scheduling decision types **inter-** and **intra-** kernel

## Halide Program for a 3x3 Convolution

### Algorithm:

```
RDom r(0, 3, 0, 3);  
output(x, y) += input(x + r.x, y + r.y)  
               * weight(r.x, r.y)
```

### Schedule:


```
input.in().store_at(output, y)  
             .compute_at(output, x);  
output.accelerate({input}, y);  
output.unroll(r.x).unroll(r.y);
```

Loop tiling,  
ordering, fusion

Which loops  
to parallelize

Memory  
hierarchy

Scope of  
accelerator



# Virtualization Summary and Next Steps

## Summary

- Efficient overlap of data movement and acceleration
- Simultaneous Applications increase accelerator utilization

## Next Steps

- Clock gate Global Buffer Unused Tiles
- Given the variety of scheduling decisions due to the complexity of our system (DMA, GLB, CGRA, etc.)
  - Tiling at DRAM level
  - Tiling at Global Buffer level
  - Overlapping of Data Movement, Kernels, Configuration etc.
- Dynamic- and Auto- Scheduling