

Amber SoC Tapeout Overview


Stanford University -- AHA Monthly Meeting

December 3, 2020

Christopher Torng, Taeyoung Kong

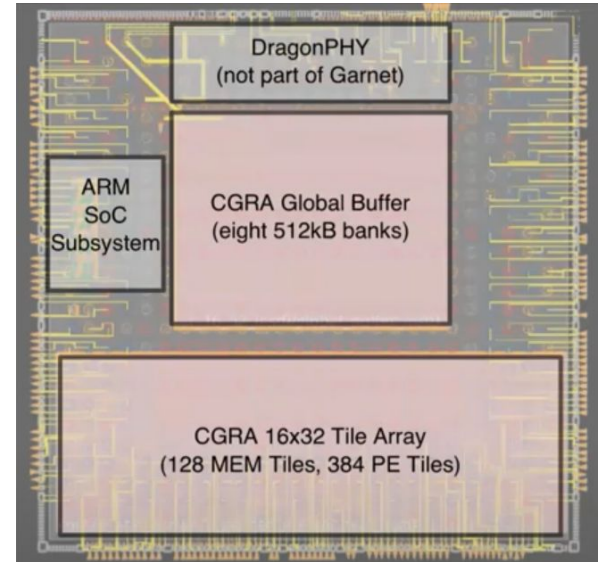
AHA SoC Evolution -- Jade CGRA

- 16x16 array of PE and memory tiles
- Simple 16b fixed point PE tiles
- Memory tiles with line buffer access pattern only
- Image processing applications



AHA SoC Evolution -- Garnet SoC


- Virtually taped out on June 2020
- Jade PE + **FP multiply and add (BFloat16)**
- Memory tiles + **general affine access pattern**
- Image processing + ML



AHA SoC Evolution -- Amber SoC

- Taped out on December 1, 2020
- Garnet PE + **Local register file (ponds)**
- Garnet memory tiles + **Easier to program from the compiler**
- Verified fast reconfiguration, application partitioning, temporal and spatial multiplexing of applications
- Image processing and ML




TSMC 16nm
Cadence VDE
5mm x 5mm



Amber SoC Architecture


Memory System for Virtualization and Energy-Efficiency

Tape-out Revisions and Memory System


Tapeout Revision	Hierarchy	Features
Jade (Gen-1 CGRA) 	Dram -> Memory Tile	First version. No SoC
Garnet (Gen-2 CGRA + SoC) 	Dram -> Global Buffer -> Memory Tile	Energy-efficiency with 3-level memory system
Amber SoC (Gen-3 CGRA + SoC) 	Dram -> Global Buffer -> Memory Tile -> Pond (Compute-local RF)	Virtualization Energy-efficiency with 4-level memory system

Introduction to CGRA Virtualization

- CGRA virtualization at different levels



Block-level



CGRA-level

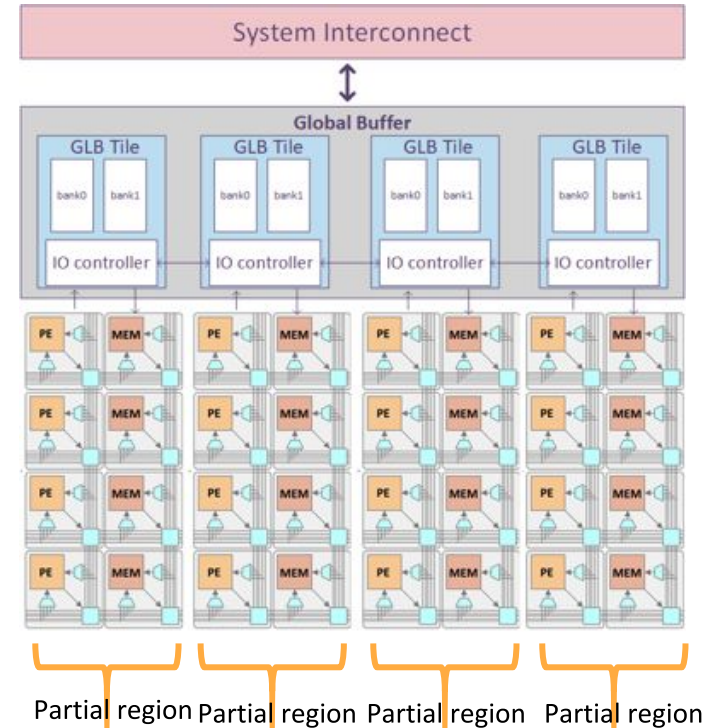


Multi node level

- Definition / Advantages
 - Abstract CGRA hardware resources
 - Simplify the interface
 - Support flexible execution models
- => Requires hardware and software support

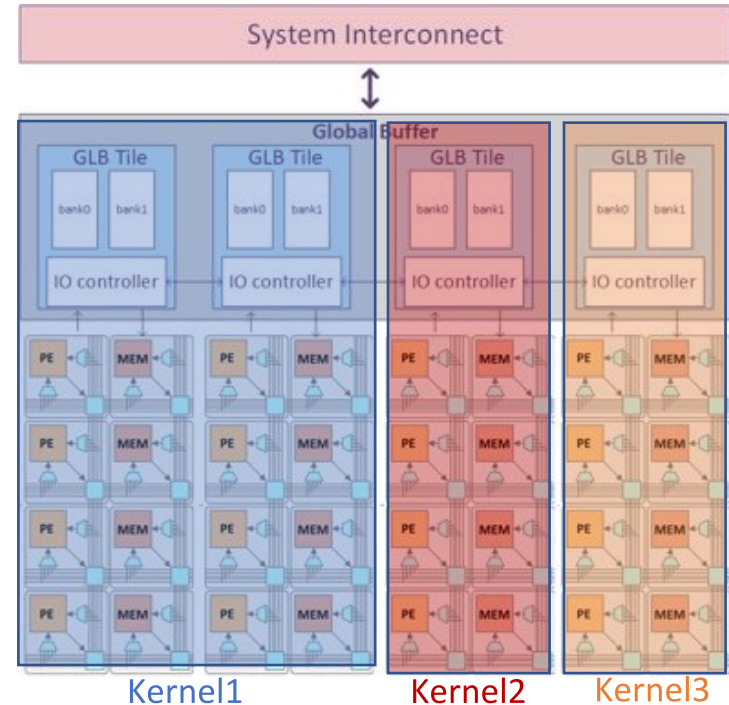
Amber: An Initial Prototype with Hardware Mechanisms for Virtualization

- Mechanism #1 : Run multiple kernels in parallel
- Mechanism #2: Relocatable bitstream to different partial regions
- Mechanism #3: Dynamic partial reconfiguration




Hardware Mechanism #1 – Multiple Kernels in Parallel

- Three kernels running independently and simultaneously
 - Virtualization at I/O, GLB, and Tile
- I/O allocation
 - Not limited to the location of the memory bank
- memory/computation allocation
 - MEM: Different # of global buffer
 - COMP: Different # of partial regions

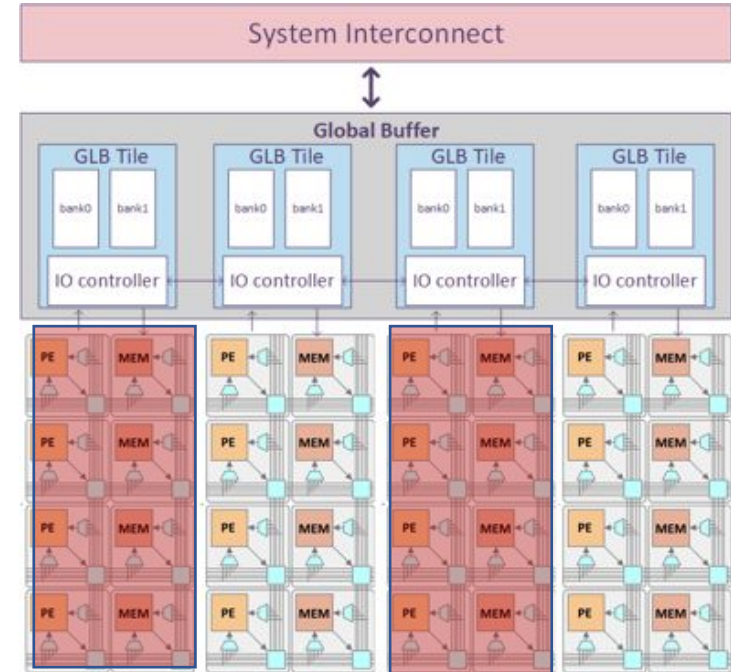


Hardware Mechanism #2 – Relocatable Bitstream to Different Partial Regions

- Bitstream address breakdown



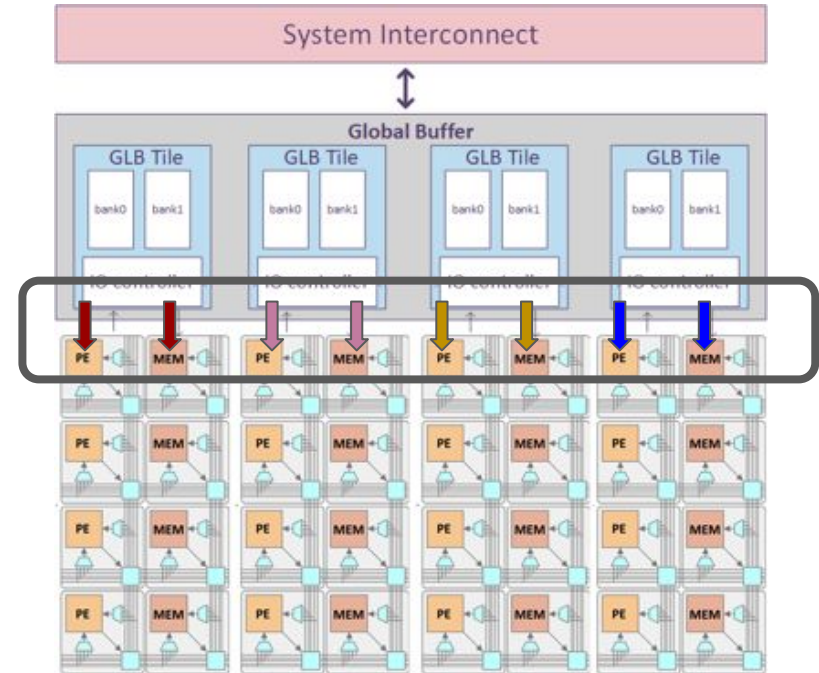
- LSB 8bits indicate the column #
- Rest address bits and data bits remain unchanged






Relocatable

Hardware Mechanism #3 – Dynamic Partial Reconfiguration (DPR)

- Use global buffer to double buffer bitstream
 - Fast
 - Parallel
 - Independent
- Configuration time (75,520 registers)
 - Jtag: 108.7ms
 - AXI4-Lite: 604us
 - DPR: 4.72us

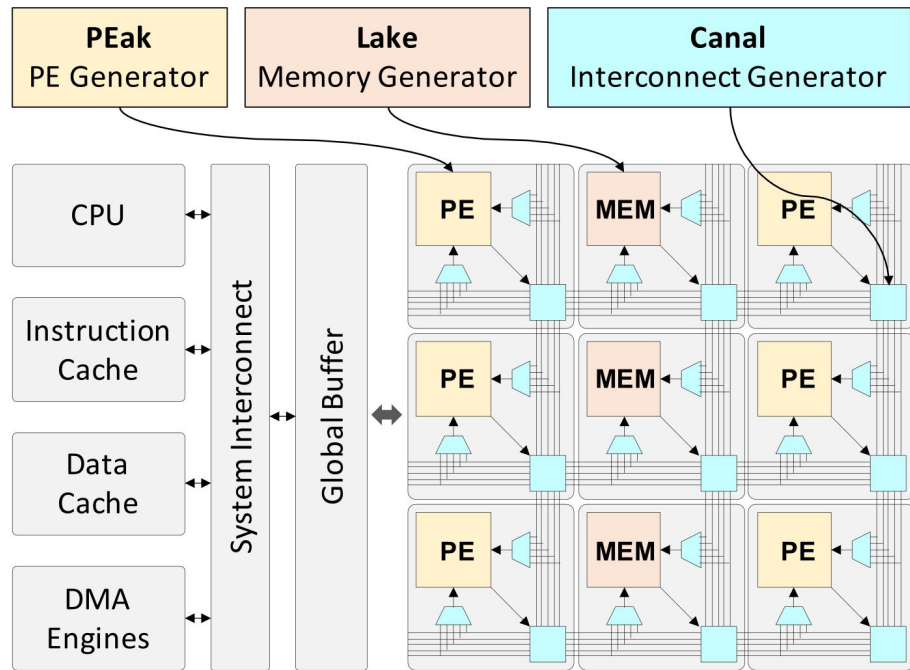


Tape-out Revisions and Memory System

Tapeout Revision		Hierarchy	Features
Jade (Gen-1 CGRA)		Dram -> Memory Tile	First version. No SoC
Garnet (Gen-2 CGRA + SoC)		Dram -> Global Buffer -> Memory Tile	Energy-efficiency with 3-level memory system
Amber SoC (Gen-3 CGRA + SoC)		Dram -> Global Buffer -> Memory Tile -> Pond (Compute-local RF)	Virtualization Energy-efficiency with 4-level memory system

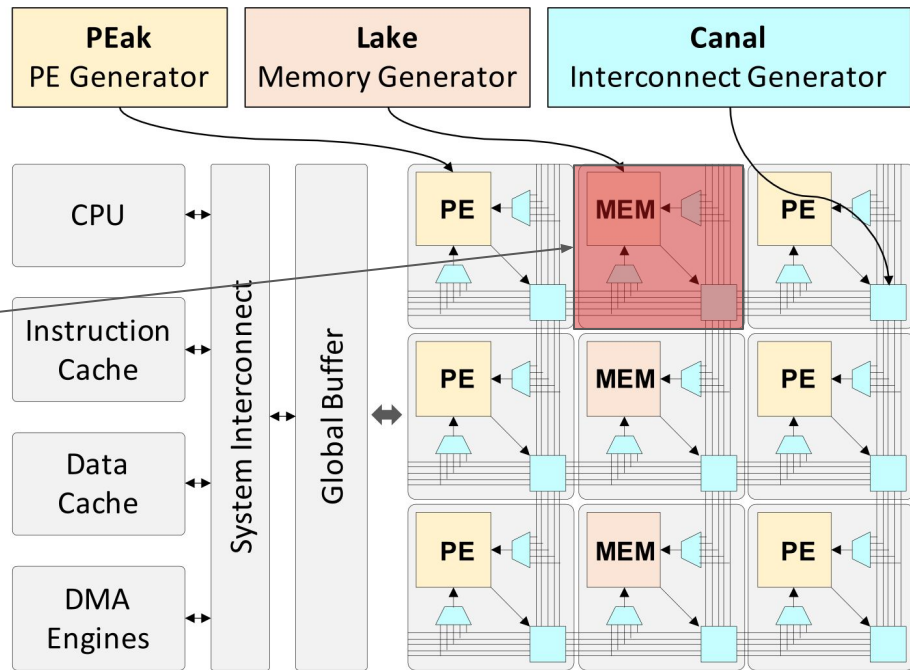
Hardware Support for Energy-Efficient Memory System

- More energy-efficient memory system
 - Reduce memory access
 - Memory hierarchy
 - Unified framework for each level of hierarchy



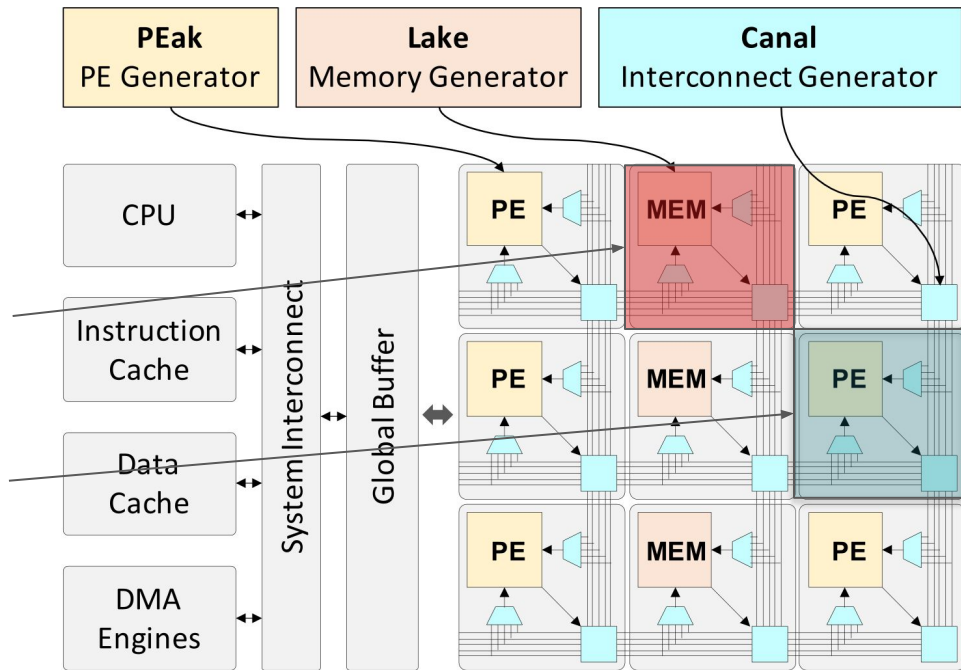
Hardware Support for Energy-Efficient Memory System

- More energy-efficient memory system
 - Reduce memory access
 - Memory hierarchy
 - Unified framework for each level of hierarchy
- Mechanism 1: Leverage wide-fetch, single-port SRAM for better efficiency



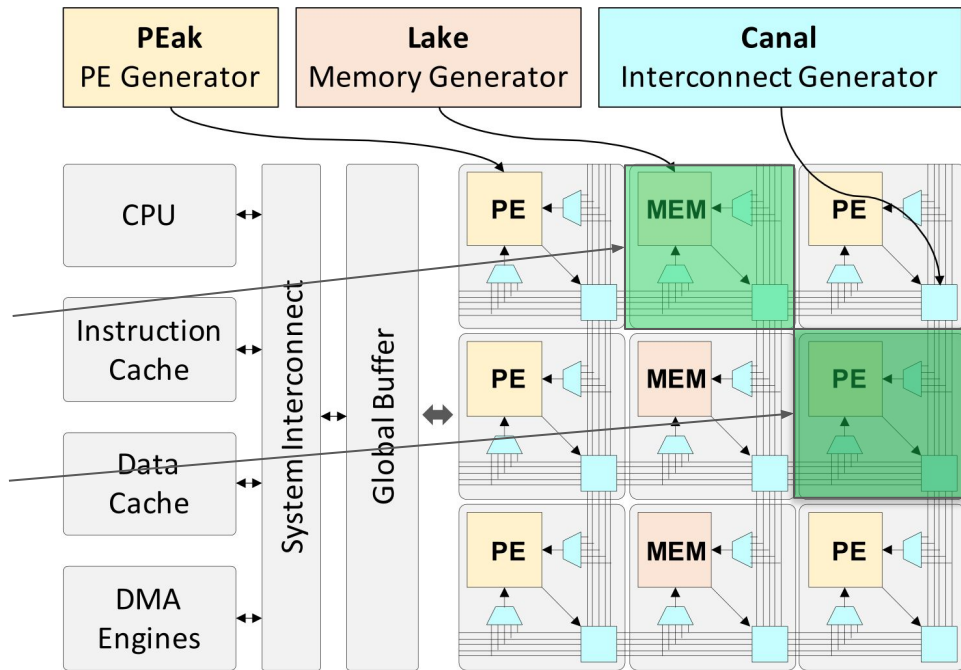
Hardware Support for Energy-Efficient Memory System

- More energy-efficient memory system
 - Reduce memory access
 - Memory hierarchy
 - Unified framework for each level of hierarchy
- Mechanism 1: Leverage wide-fetch, single-port SRAM for better efficiency
- Mechanism 2: Place small register-based memories in the PE




Hardware Support for Energy-Efficient Memory System

- More energy-efficient memory system
 - Reduce memory access
 - Memory hierarchy
 - Unified framework for each level of hierarchy
- Mechanism 1: Leverage wide-fetch, single-port SRAM for better efficiency
- Mechanism 2: Place small register-based memories in the PE
- Mechanism 3: Unified memory generator framework for every level of hierarchy




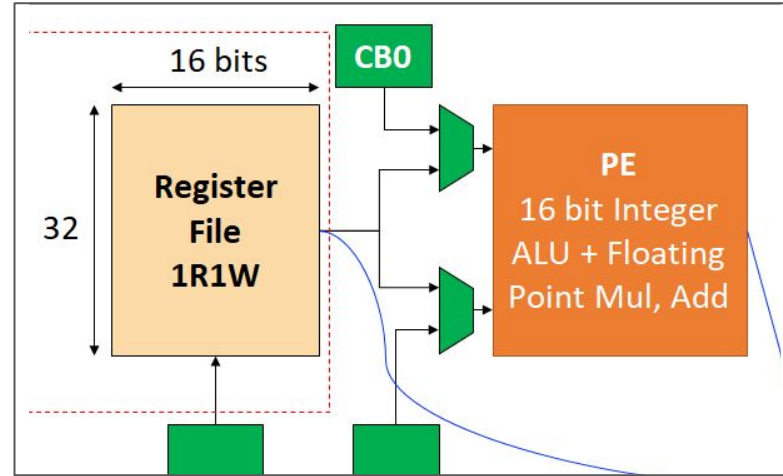
Hardware Mechanism #1 - Wide Fetch SRAM

- Mechanism 1: Leverage wide-fetch, single-port SRAM
 - Reduce memory read/write
- Wide-fetch SRAMs exhibit better energy per data accessed compared to single-wide SRAMs
- Wide-fetch SRAMs provide equivalent bandwidth with only a modest control overhead compared to Dual-Ported SRAMs.




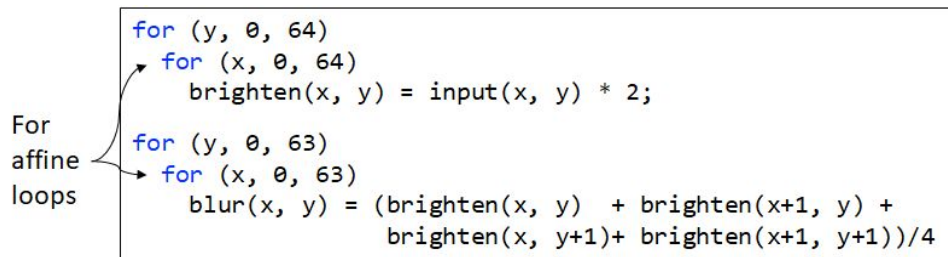
Hardware Mechanism #2 - PE-local Register File

- Mechanism 2: Register file in the PE
 - The lowest level memory hierarchy
- Keep data local to PE for accumulation and low-latency buffering
- Avoid SRAM accesses by exploiting locality



Hardware Mechanism #3 - Unified Framework

So far, the discussion has concerned mainly architecture...what about control?




- Mechanism 3: Unified framework of memory generator
 - Use the same framework for every memory level in the hierarchy
 - Different SRAM organization and addressing complexity
- Compiler analyzes looping structures of the application with given resource constraints
 - Generate configuration registers

Summary

- **Amber** is a prototype for hardware mechanisms for CGRA **virtualization**
 - Multi-kernels in parallel
 - Relocatable bitstream
 - Dynamic partial reconfiguration
- **Amber** has an **energy-efficient** multi-level memory system
 - Wide-fetch, single-port SRAM in MEM tile
 - Register-file in PE tile
 - Unified memory generator framework for each level of memory

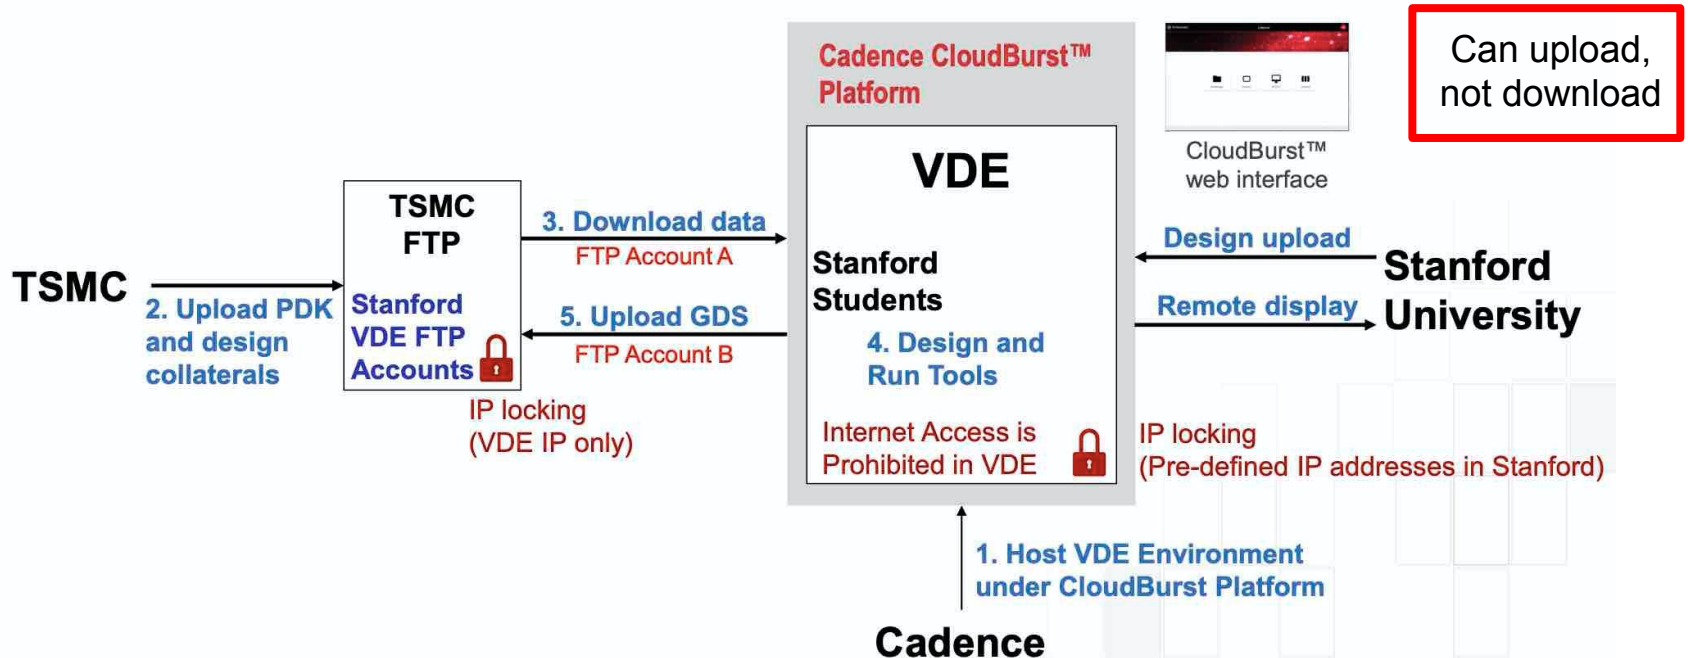
Tapeout Methodology

1. Building a chip in the cloud
2. Reflecting on a first attempt at agile physical development



We built Amber in TSMC Virtual Design Environment (VDE)

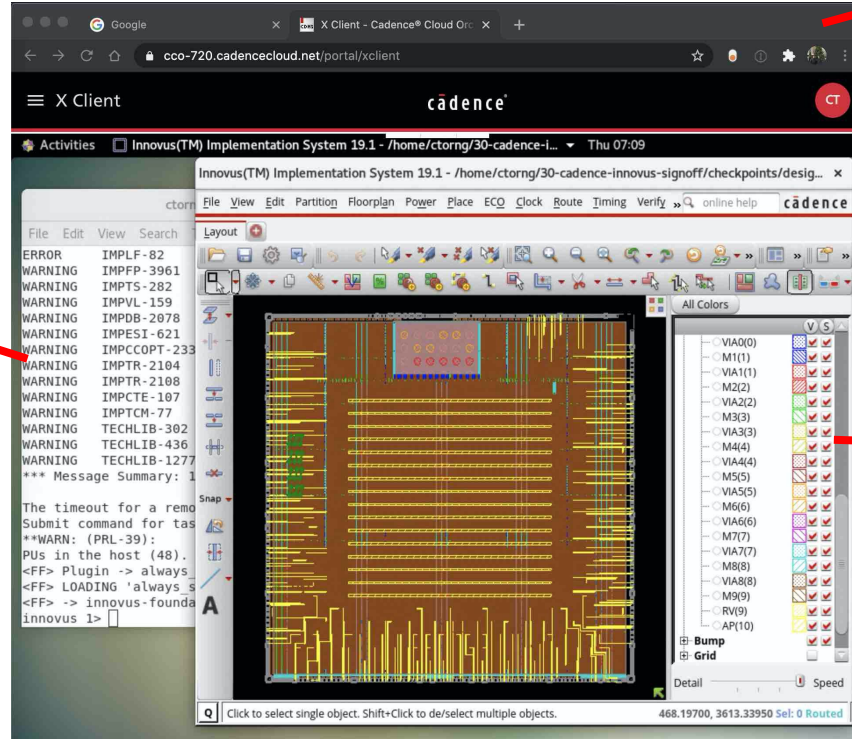
- **Vendor:** Prevent accidental leakage of TSMC's technology information
- **Customer:** Cloud EDA



We built Amber in TSMC Virtual Design Environment (VDE)

Access to terminal

Logging in from a browser



Display available for GUI needs

VDE was a success in many respects

From TSMC and Cadence's perspective

- TSMC protected their sensitive information from accidental leakage
- Cadence has a new platform to offer as a service

From Stanford's perspective


- We **successfully built a chip** in the cloud
- We had access to all the tools, IP, and compute we needed
- Unique for academics -- We had access to live EDA support (Cadence)

VDE introduced new pain points

Challenge	Why it can hurt
Not allowed to extract design and source files from the closed environment (or difficult to)	Losing project history hurts, especially for open-source projects. Difficult to share code with non-VDE projects.
Cloud platform sometimes high latency	Interactive jobs can be very slow
Each VDE instance is a “new system”	Need to spend time getting tools to work again
Single-vendor toolchain (Cadence)	Less flexibility + Team may not have experts
Having to juggle cloud resources vs costs (e.g., making sure big jobs run on big servers)	Overpay for cloud resources or suffer the low performance

Our take on the VDE

- Modularity with **mflowgen** helped us quickly adapt between a single-vendor flow (Cadence) and a mixed-vendor flow (Synopsys, Cadence, Mentor)



- We maintained version control outside the VDE and **repeatedly uploaded our changes** (only lightweight changes)
- Other challenges decreased productivity (robustness, latency, tools, maintenance), but the **chip still went out successfully**


Tapeout Methodology

1. Building a chip in the cloud

- VDE was a success with room for improvement

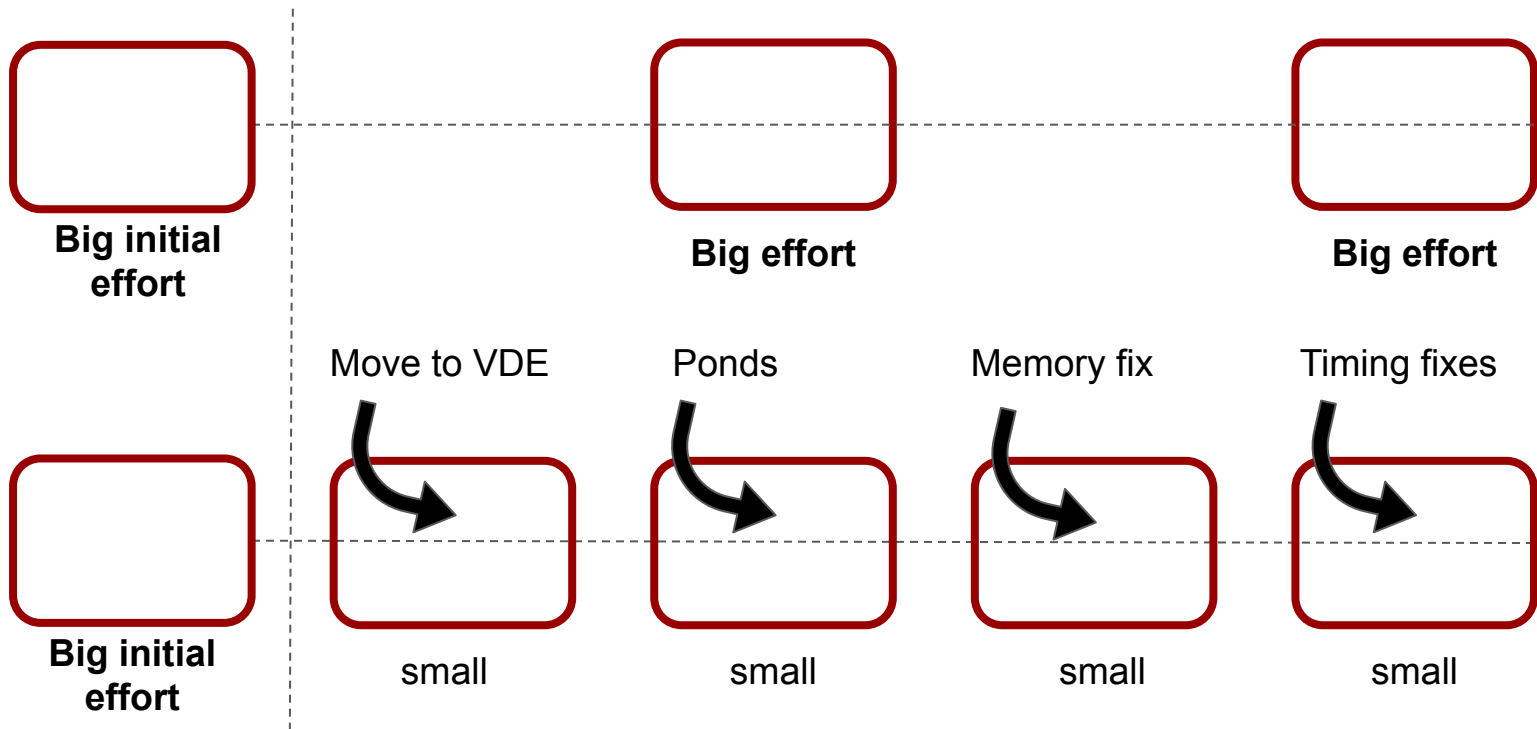
2. Reflecting on a first attempt at agile physical development

- Producing a golden tape every week




Weekly golden releases instead of “one big release”

Push away from “**one big release**” and towards software-like **frequent small ones**



Weekly golden releases instead of “one big release”

- We were successful to a large degree...




```
% ls gold.archives  
  
gold.112  
gold.148  
gold.158  
gold.174  
gold.1987  
gold.30  
gold.40  
gold.84  
gold.93  
gold.README
```

A history of golden builds
throughout design
development

- Required **culture change** and also surfaced **fundamental challenges**

Agile culture shift -- Biweekly sprints for team focus



Biweekly sprints helped us convert many unordered parallel efforts into a **progression of team efforts** that “commit” one by one into the chip

Agile culture shift -- Adapting Scrums for COVID

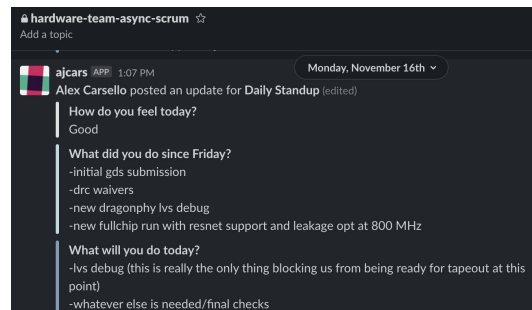


Geekbot

Questions

- How do you feel today?
- What did you do since {last_report_date}?
- What will you do today?
- Anything blocking your progress?

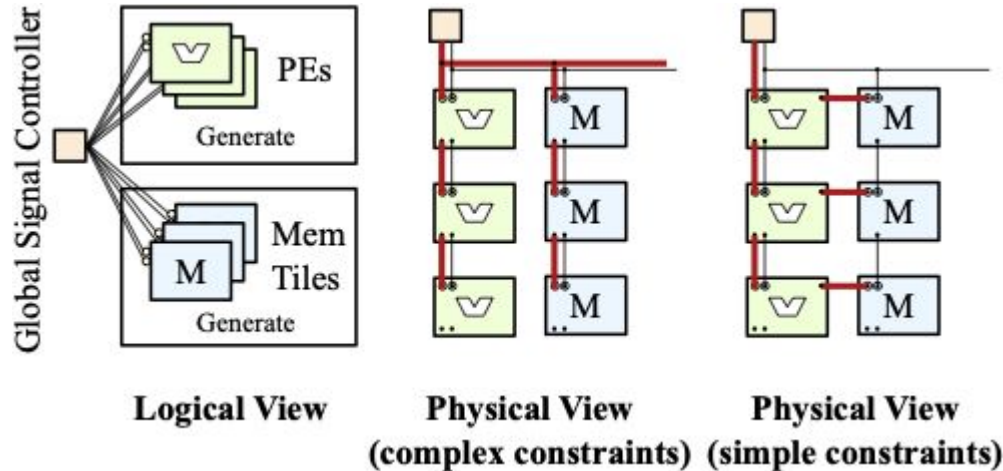
Participants



Slack channel
with everyone's
responses

Asynchronous scrums completely replaced synchronous ones to great success (e.g., staying in sync, text records, less meeting fatigue)

Challenge #1 -- Some design approaches are friendlier than others




The (left) physical view requires **manual attention to constraints**.

The (right) physical view requires **negligible human attention**.

While difficult to spot, certain design decisions which have no impact on logical design can cost weeks of physical iteration.

These need to be eliminated.

Challenge #2 -- No single source of truth between physical/logical



Significant design churn when...

- The **logical design updates**, but **physical collateral is stale** (e.g., constraints)
- The logical design updates **non-functionally** but generates a different netlist which breaks physical scripts

Need a mechanism to better tie together related logical and physical design files

Can we address these two challenges?

- Challenge #1 -- Some design approaches are friendlier than others
- Challenge #2 -- No single source of truth between physical/logical

Summary of Amber SoC tapeout

- Taped out on December 1, 2020
- Garnet PE + **Local register file (ponds)**
- Garnet memory tiles + **Easier to program from the compiler**
- Verified fast reconfiguration, application partitioning, temporal and spatial multiplexing of applications
- Lessons learned in the **cloud**
- Lessons learned for **agile design**

TSMC 16nm
Cadence VDE
5mm x 5mm

